

# Deep Learning- CSE641

## Assignment - 3

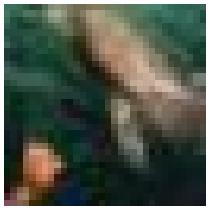
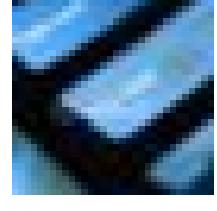
**Name:** Akanksha Shrimal  
**Name:** Vaibhav Goswami  
**Name:** Shivam Sharma

**Roll No:** MT20055  
**Roll No:** MT20018  
**Roll No:** MT20121

FILES SUBMITTED : submitted .py file , .ipynb file and readme.pdf and output.pdf

## 1. Convolutional Neural Network (CNN):

### 1.1 Visualize 10 Samples From Each Class:

Class	Image 1	Image 2
0		
1		
2		
3		
4		
5		

6		
7		
8		
9		

## 1.2 CNN Architecture

### 1.2.1 Activation function: Tanh

- **Block 1**

```

Model_B1(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): Tanh()
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): Tanh()
        (6): MaxPool2d(kernel_size=4, stride=3, padding=0, dilation=1,
ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=1600, out_features=512, bias=True)
    )
)

```

```

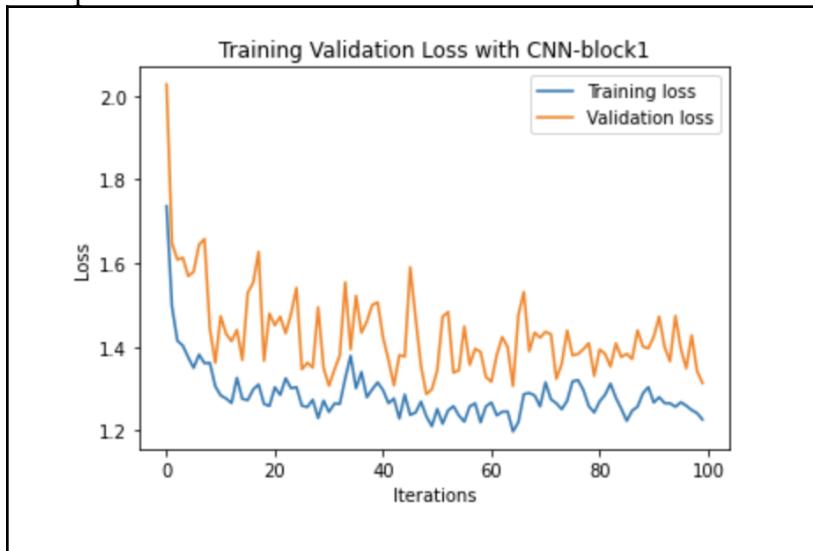
(1): Tanh()
(2): Linear(in_features=512, out_features=256, bias=True)
(3): Tanh()
(4): Linear(in_features=256, out_features=64, bias=True)
(5): Tanh()
(6): Linear(in_features=64, out_features=10, bias=True)
)
)

```

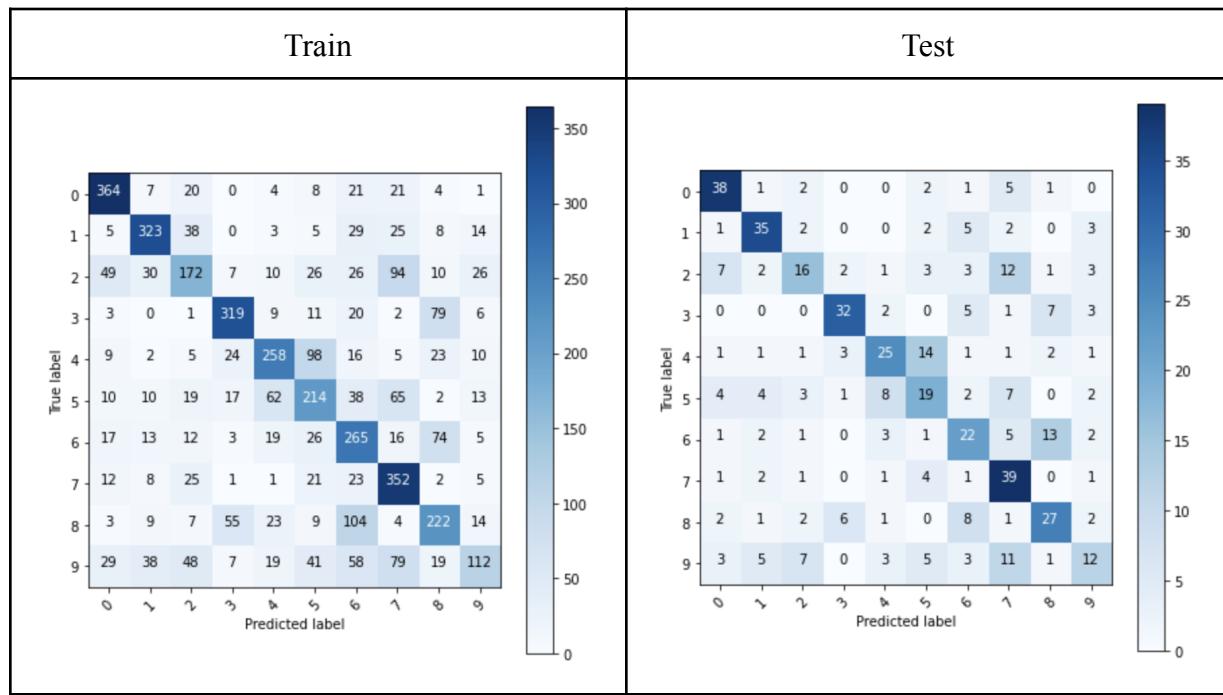
❖ Accuracy and Loss

Block 1	Train	Test
Accuracy	57%	53%
Loss	1.200912	1.292314

❖ Loss plot



## ❖ Confusion Matrix



### ● Block 1,2

```

Model_B12(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): Tanh()
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): Tanh()
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (9): Tanh()
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): Tanh()
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): Tanh()
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=512, out_features=256, bias=True)
        (1): Tanh()
    )
)

```

```

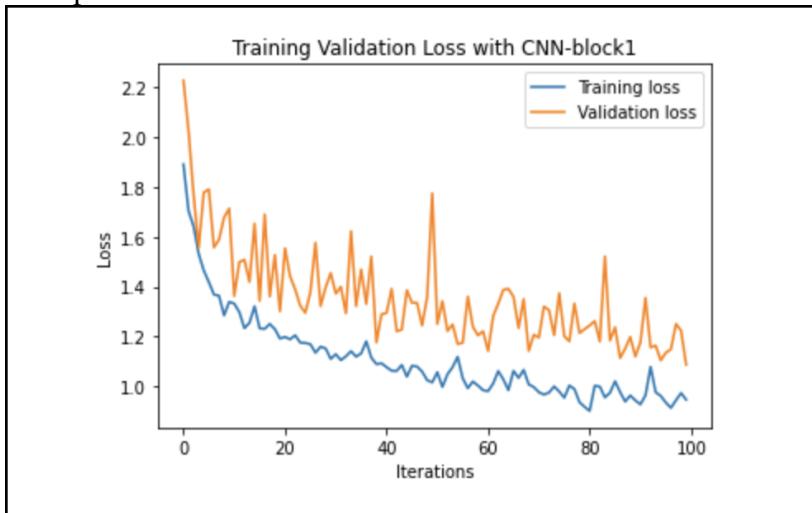
(2) : Linear(in_features=256, out_features=64, bias=True)
(3) : Tanh()
(4) : Linear(in_features=64, out_features=10, bias=True)
)
)

```

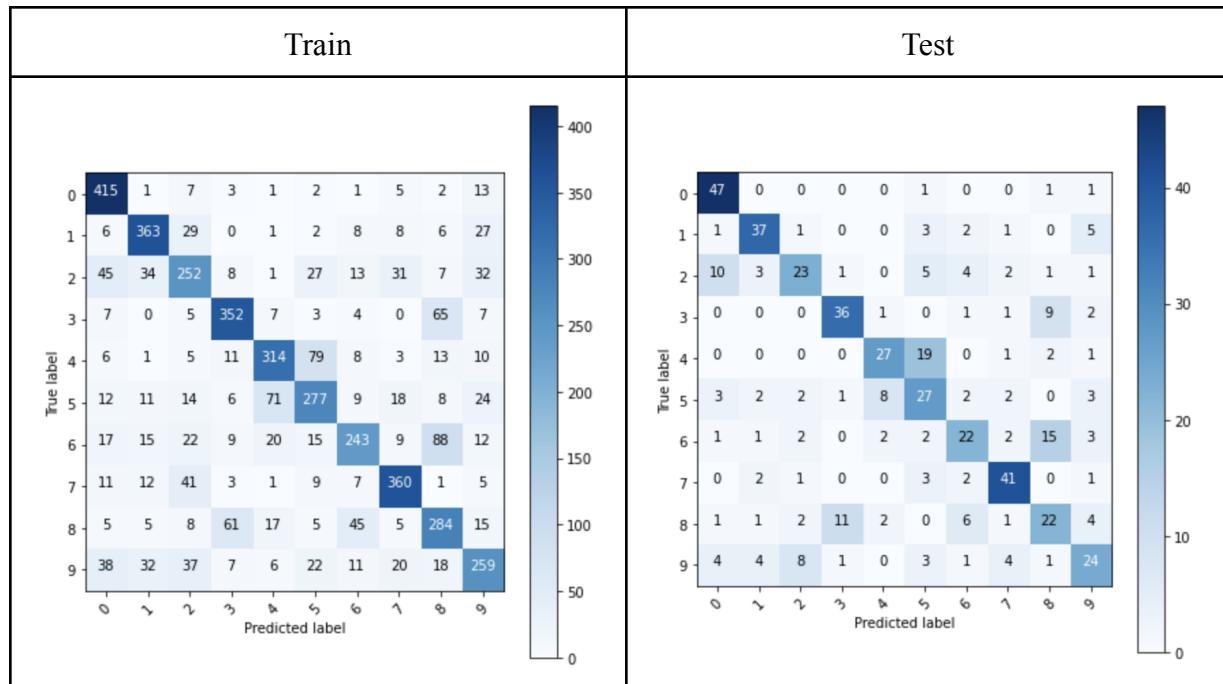
❖ Accuracy and Loss

Block 1,2	Train	Test
Accuracy	69%	61%
Loss	0.894667	1.89120

❖ Loss plot



❖ Confusion Matrix



- Block 1,2,3

```

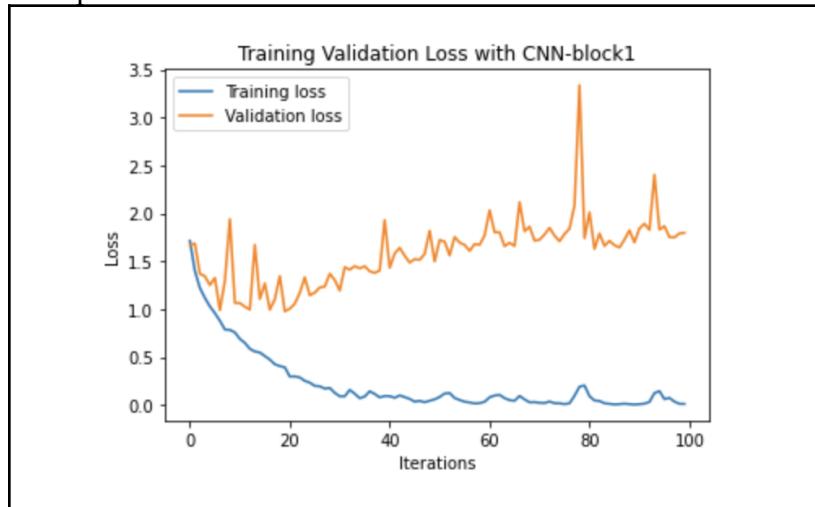
Model_B123(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): Tanh()
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): Tanh()
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (9): Tanh()
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): Tanh()
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): Tanh()
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (19): Tanh()
        (20): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)

```

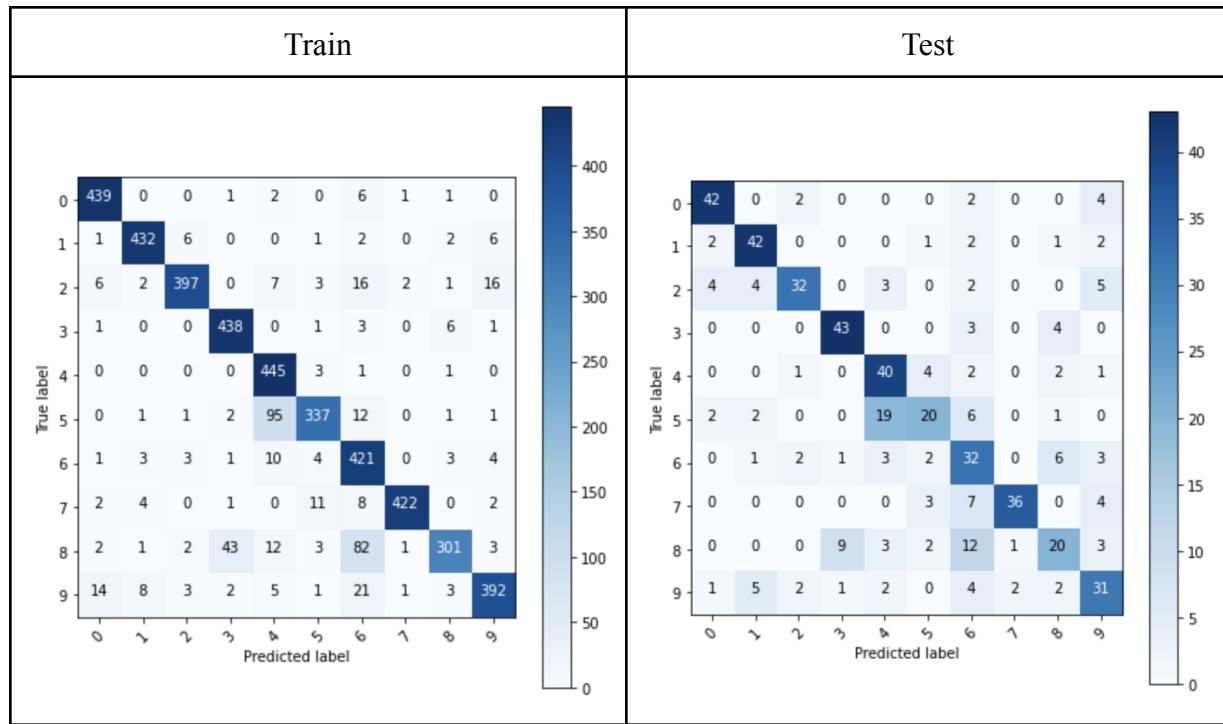
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	89%	67%
Loss	0.301188	0.963042

❖ Loss plot



❖ Confusion Matrix



### 1.2.2 Activation function: ReLU

- Block 1

```
Model_B1(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
```

```

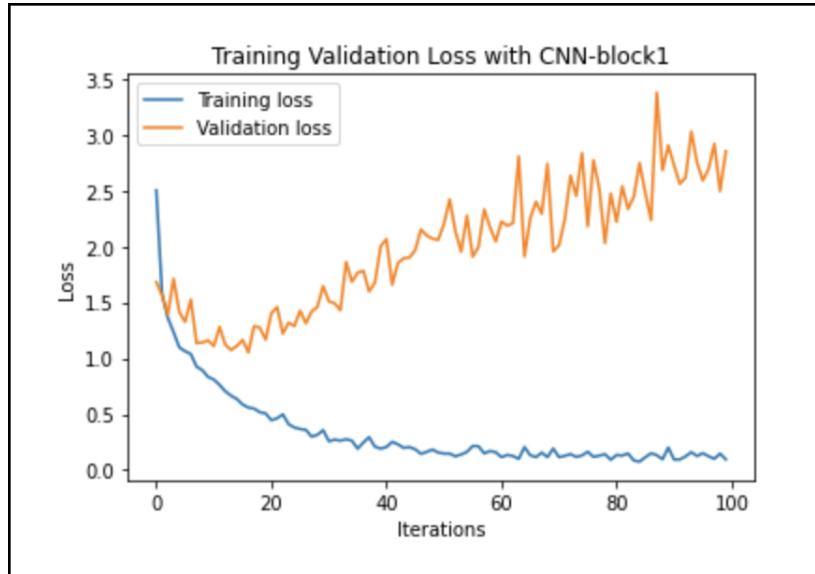
        track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=4, stride=3, padding=0, dilation=1,
ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=1600, out_features=512, bias=True)
        (1): ReLU(inplace=True)
        (2): Linear(in_features=512, out_features=256, bias=True)
        (3): ReLU(inplace=True)
        (4): Linear(in_features=256, out_features=64, bias=True)
        (5): ReLU(inplace=True)
        (6): Linear(in_features=64, out_features=10, bias=True)
    )
)

```

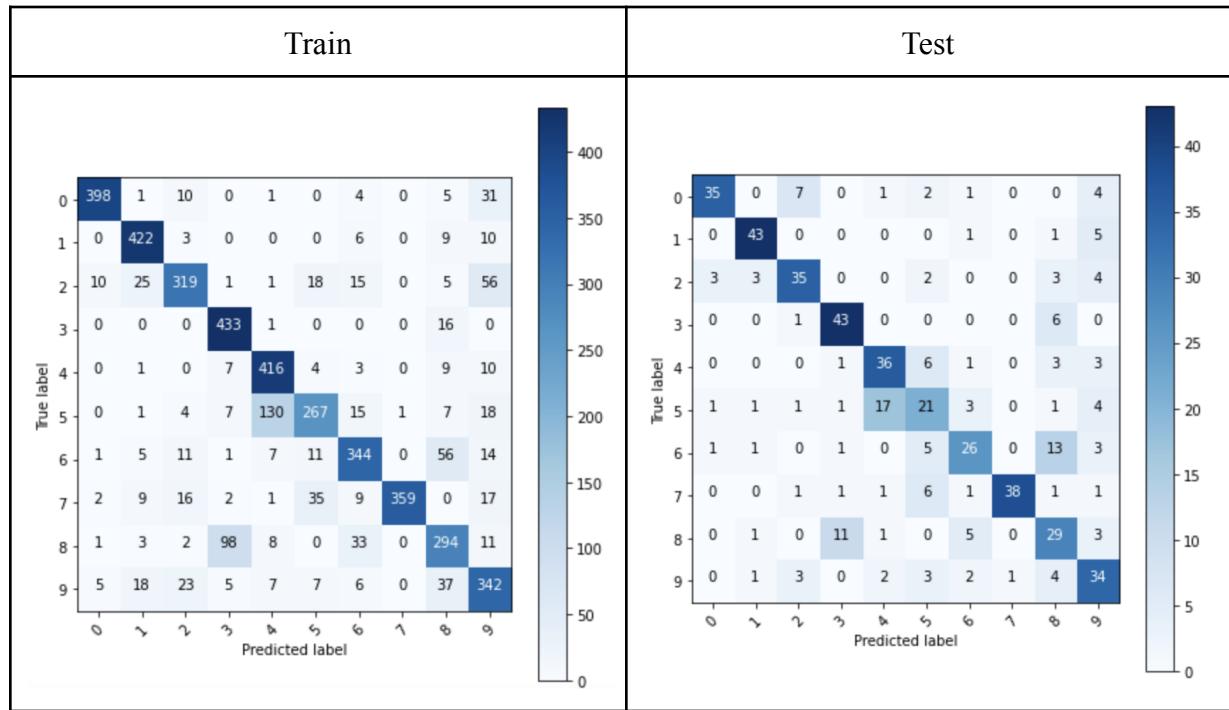
❖ Accuracy and Loss

Block 1	Train	Test
Accuracy	79%	68%
Loss	0.549036	1.057941

❖ Loss plot



## ❖ Confusion Matrix



- Block 1,2

```

Model_B12(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (9): ReLU(inplace=True)
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): ReLU(inplace=True)
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    )
    classifier): Sequential(
        (0): Linear(in_features=512, out_features=256, bias=True)
    )
)

```

```

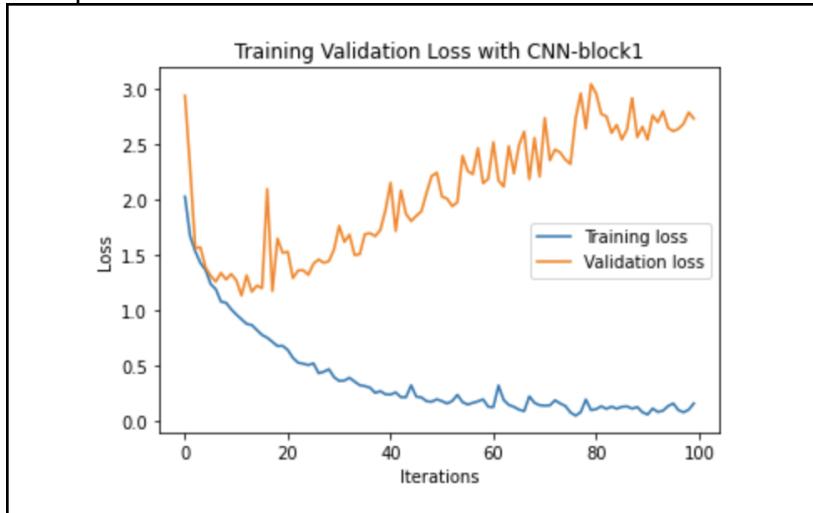
(1) : ReLU(inplace=True)
(2) : Linear(in_features=256, out_features=64, bias=True)
(3) : ReLU(inplace=True)
(4) : Linear(in_features=64, out_features=10, bias=True)
)
)

```

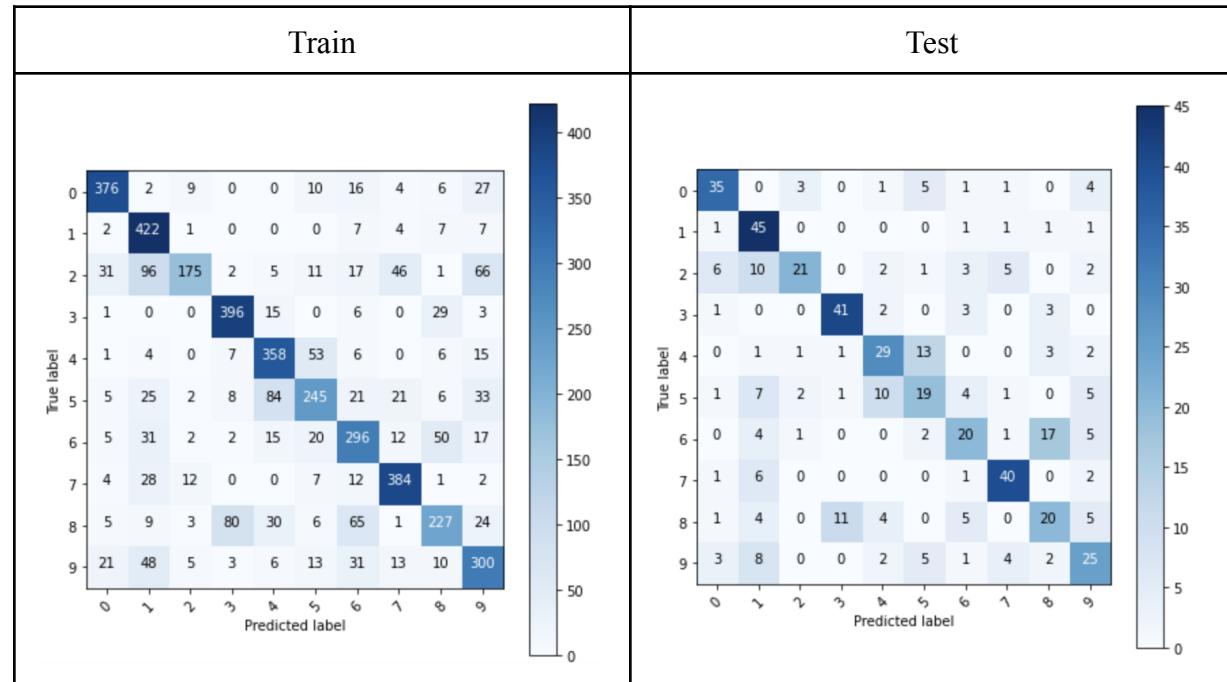
❖ Accuracy and Loss

Block 1,2	Train	Test
Accuracy	70%	59%
Loss	0.837771	1.137523

❖ Loss plot



❖ Confusion Matrix



- Block 1,2,3

```

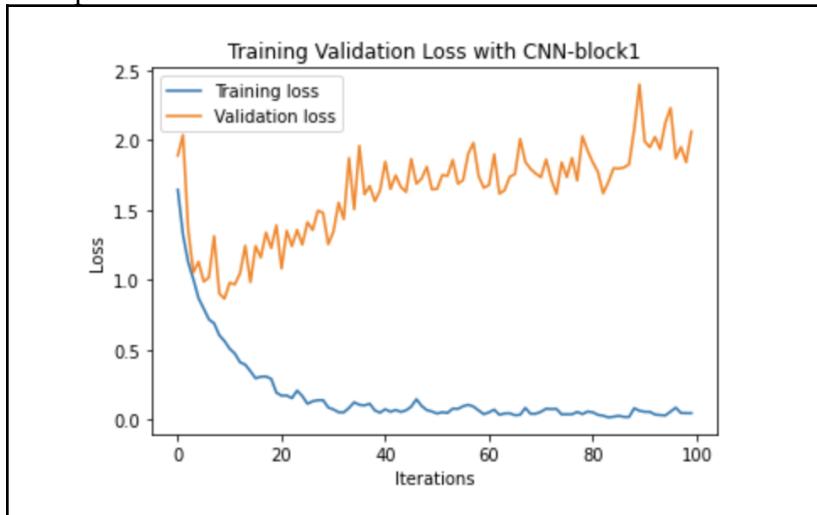
Model_B123(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (9): ReLU(inplace=True)
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): ReLU(inplace=True)
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (19): ReLU(inplace=True)
        (20): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)

```

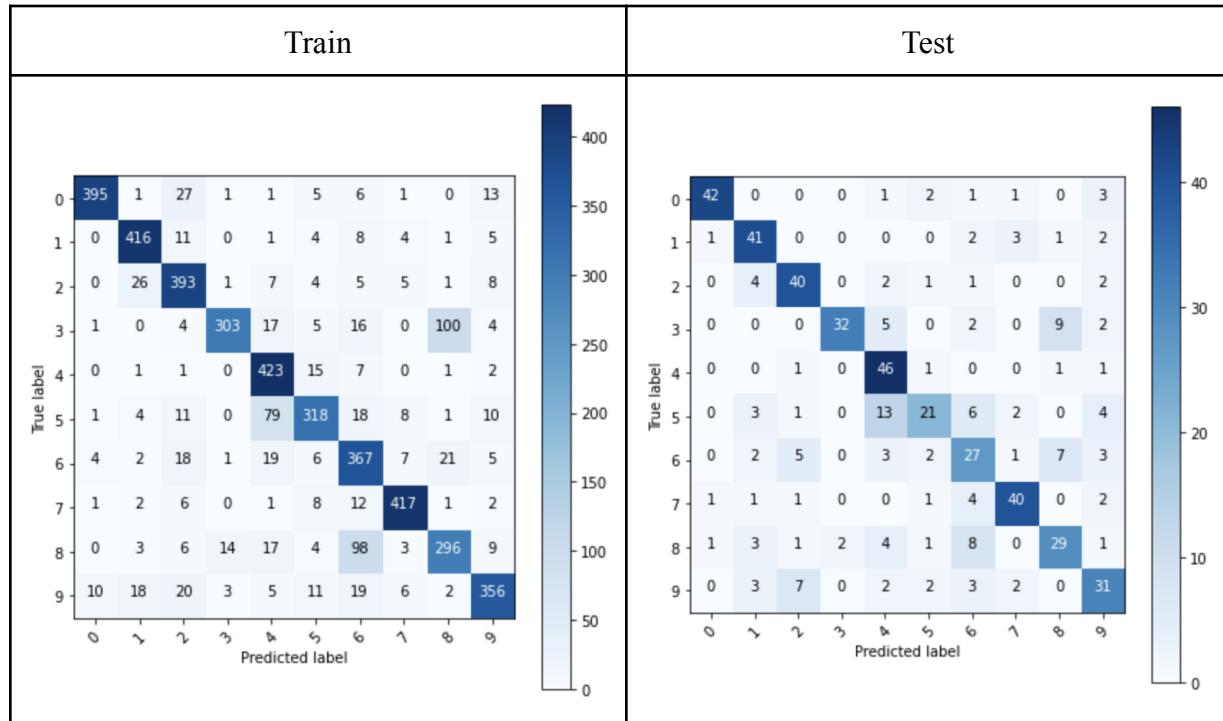
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	81%	69%
Loss	0.512706	0.868855

❖ Loss plot



❖ Confusion Matrix



**Analysis:**

- Blocks were added one at a time and performance was evaluated for each model.
- With the addition of blocks, the performance increased, i.e. accuracy increased and loss decreased in both cases (tanh and ReLU).
- ReLU performed better on average than tanh. Therefore, for the rest of the question, ReLU activation function is used.

### 1.3 Dropout Probability

#### 1.3.1 Dropout after Convolutional layers

### 1.3.1.1 Dropout = 0.3

```

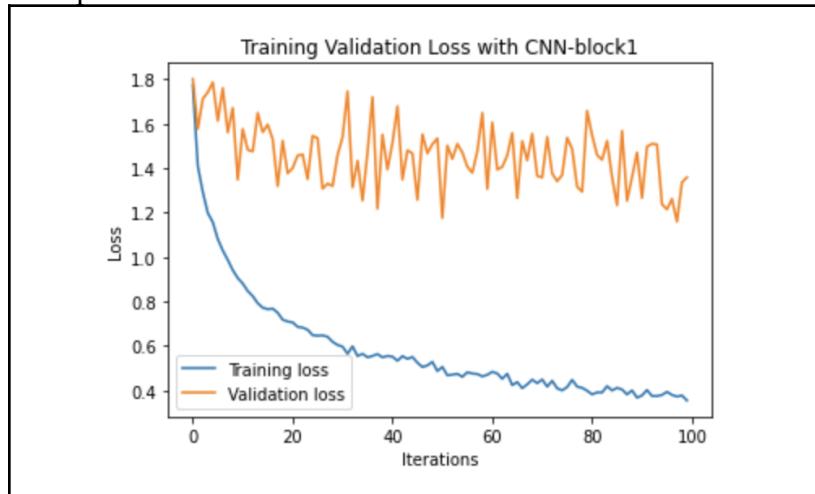
Model_B123(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Dropout(p=0.3, inplace=False)
        (4): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (5): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (6): ReLU(inplace=True)
        (7): Dropout(p=0.3, inplace=False)
        (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (9): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (10): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (11): ReLU(inplace=True)
        (12): Dropout(p=0.3, inplace=False)
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): ReLU(inplace=True)
        (16): Dropout(p=0.3, inplace=False)
        (17): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (18): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (19): ReLU(inplace=True)
        (20): Dropout(p=0.3, inplace=False)
        (21): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (22): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (23): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (24): ReLU(inplace=True)
        (25): Dropout(p=0.3, inplace=False)
        (26): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)

```

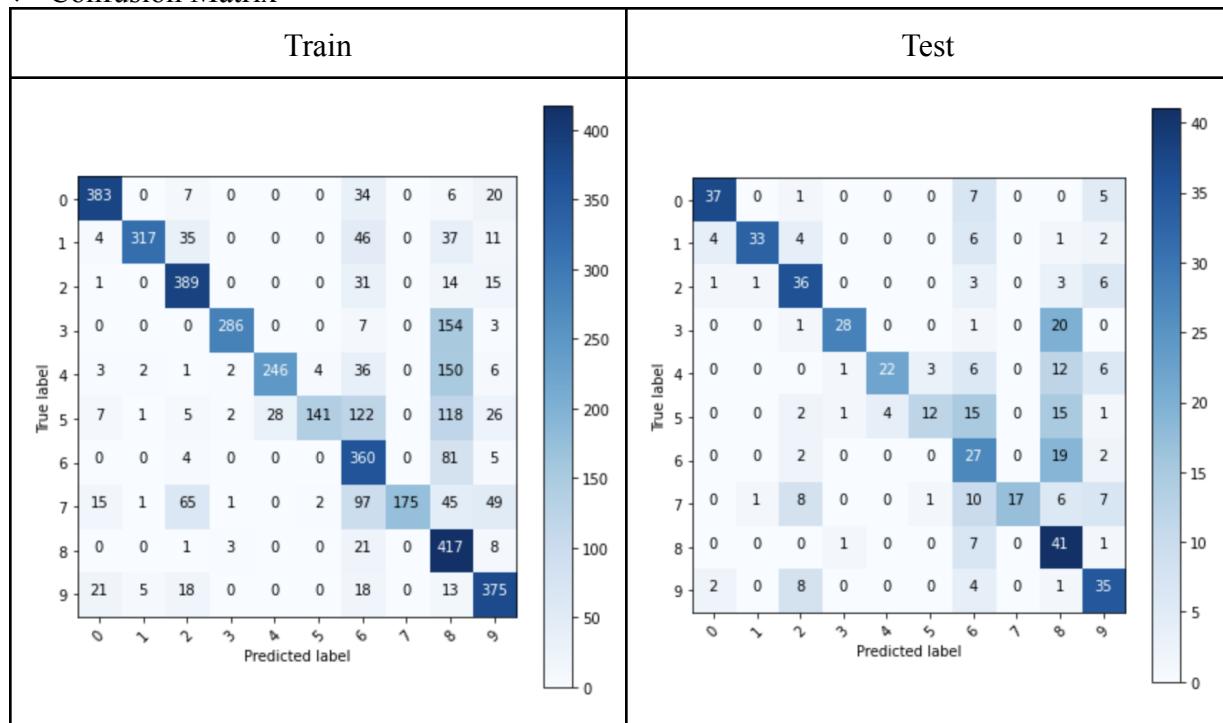
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	68%	57%
Loss	0.903236	1.165431

❖ Loss plot



❖ Confusion Matrix



### 1.3.1.2 Dropout = 0.5

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Dropout(p=0.5, inplace=False)
        (4): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (5): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (6): ReLU(inplace=True))
```

```

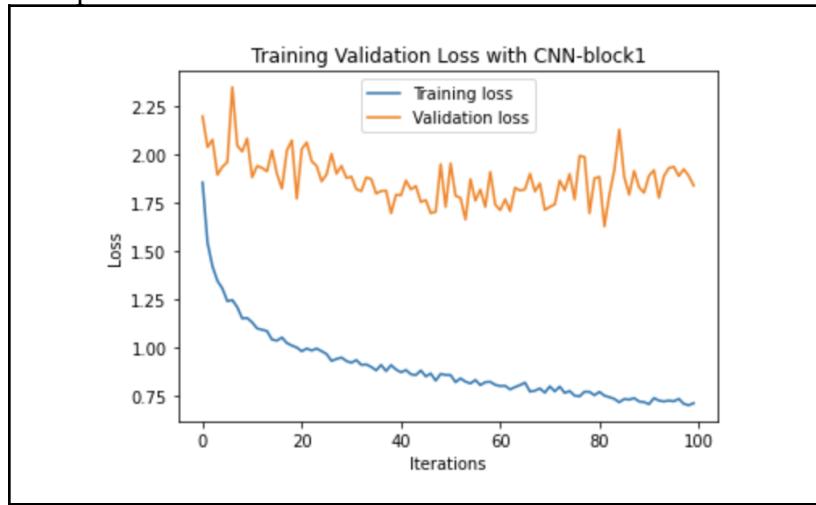
(7): Dropout(p=0.5, inplace=False)
(8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(9): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
(10): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(11): ReLU(inplace=True)
(12): Dropout(p=0.5, inplace=False)
(13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(15): ReLU(inplace=True)
(16): Dropout(p=0.5, inplace=False)
(17): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(18): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(19): ReLU(inplace=True)
(20): Dropout(p=0.5, inplace=False)
(21): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
(22): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(23): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(24): ReLU(inplace=True)
(25): Dropout(p=0.5, inplace=False)
(26): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

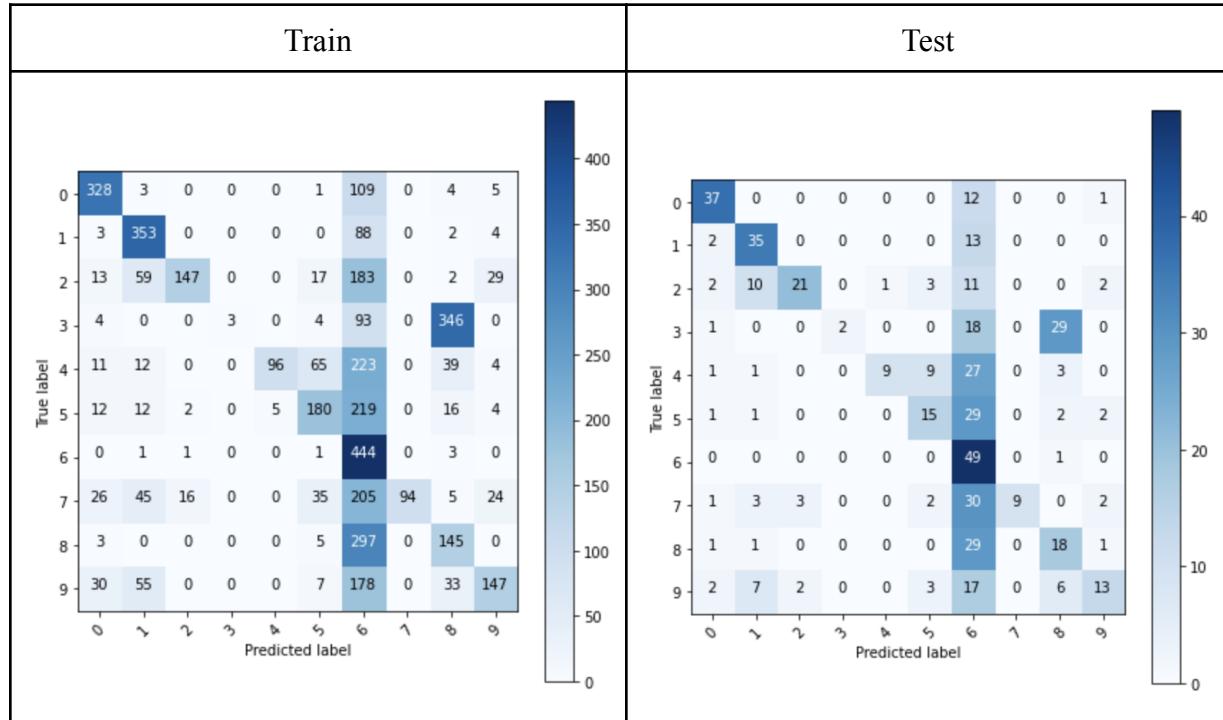
## ❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	43%	41%
Loss	1.542332	1.630709

❖ Loss plot



❖ Confusion Matrix



### 1.3.1.2 Dropout = 0.7

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Dropout(p=0.7, inplace=False)
        (4): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (5): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (6): ReLU(inplace=True)
        (7): Dropout(p=0.7, inplace=False)
        (8): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
```

```

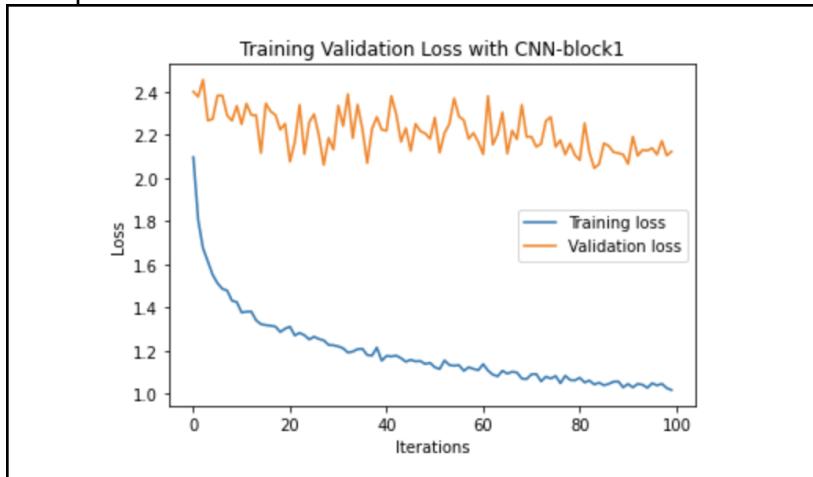
ceil_mode=False)
(9): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
(10): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(11): ReLU(inplace=True)
(12): Dropout(p=0.7, inplace=False)
(13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(15): ReLU(inplace=True)
(16): Dropout(p=0.7, inplace=False)
(17): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(18): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(19): ReLU(inplace=True)
(20): Dropout(p=0.7, inplace=False)
(21): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
(22): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(23): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(24): ReLU(inplace=True)
(25): Dropout(p=0.7, inplace=False)
(26): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Linear(in_features=64, out_features=10, bias=True)
)
)
)

```

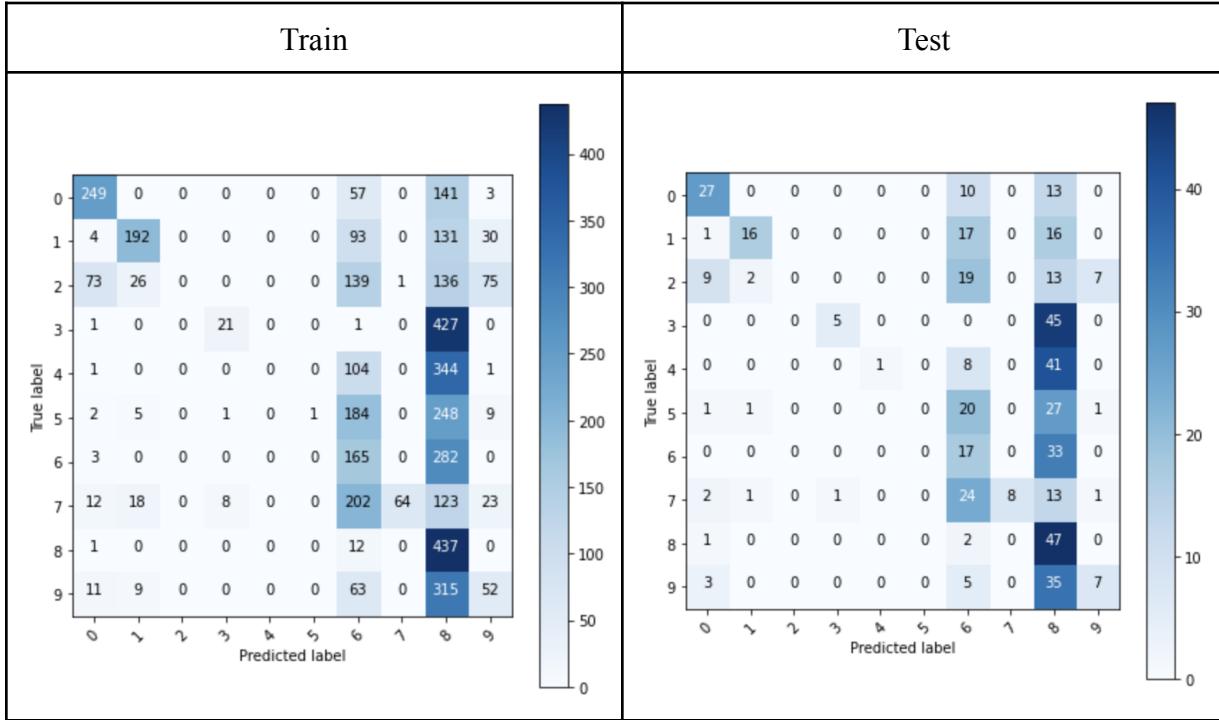
#### ❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	26%	25%
Loss	2.020246	2.043498

#### ❖ Loss plot



## ❖ Confusion Matrix



### 1.3.2 Dropout between fully-connected layers

#### 1.3.2.1 Dropout = 0.3

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (9): ReLU(inplace=True)
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): ReLU(inplace=True)
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
```

```

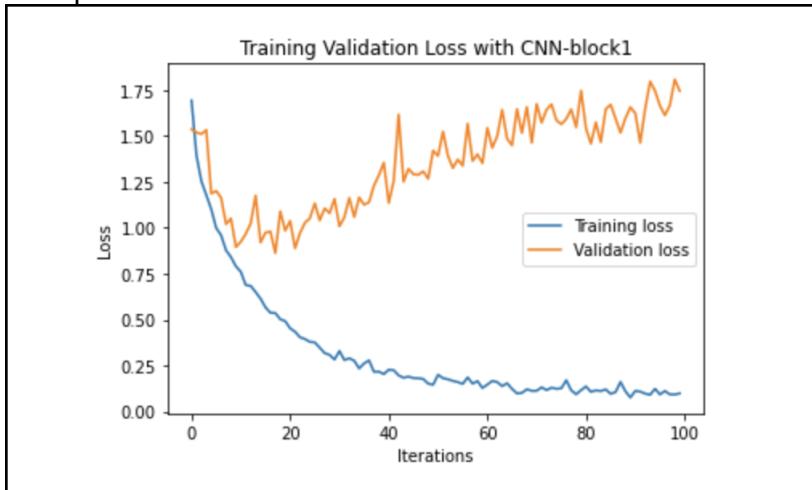
(18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(19): ReLU(inplace=True)
(20): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Dropout(p=0.3, inplace=False)
(1): Linear(in_features=64, out_features=10, bias=True)
)
)

```

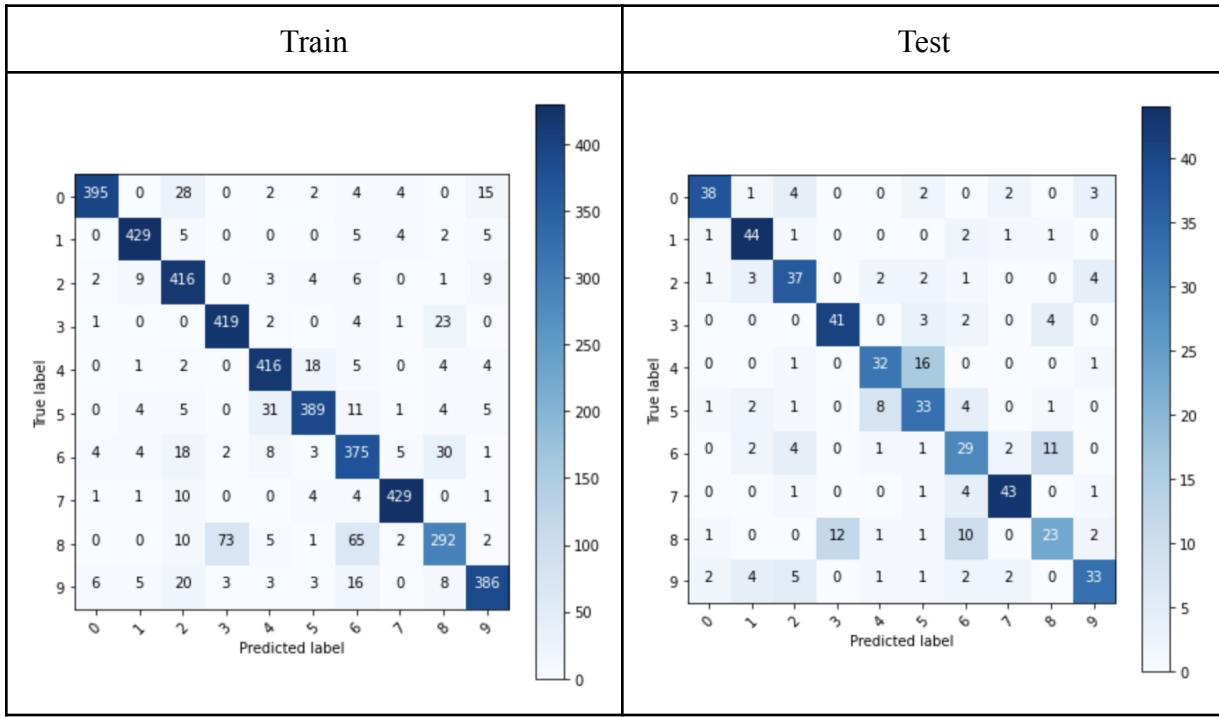
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	87%	70%
Loss	0.355316	0.854087

❖ Loss plot



## ❖ Confusion Matrix



### 1.3.2.2 Dropout = 0.5

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (9): ReLU(inplace=True)
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): ReLU(inplace=True)
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
```

```

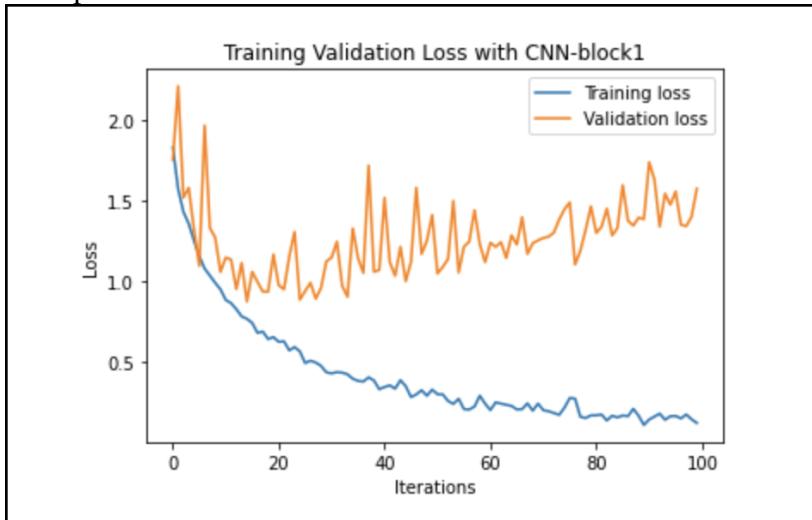
(19): ReLU(inplace=True)
(20): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Dropout(p=0.5, inplace=False)
(1): Linear(in_features=64, out_features=10, bias=True)
)
)

```

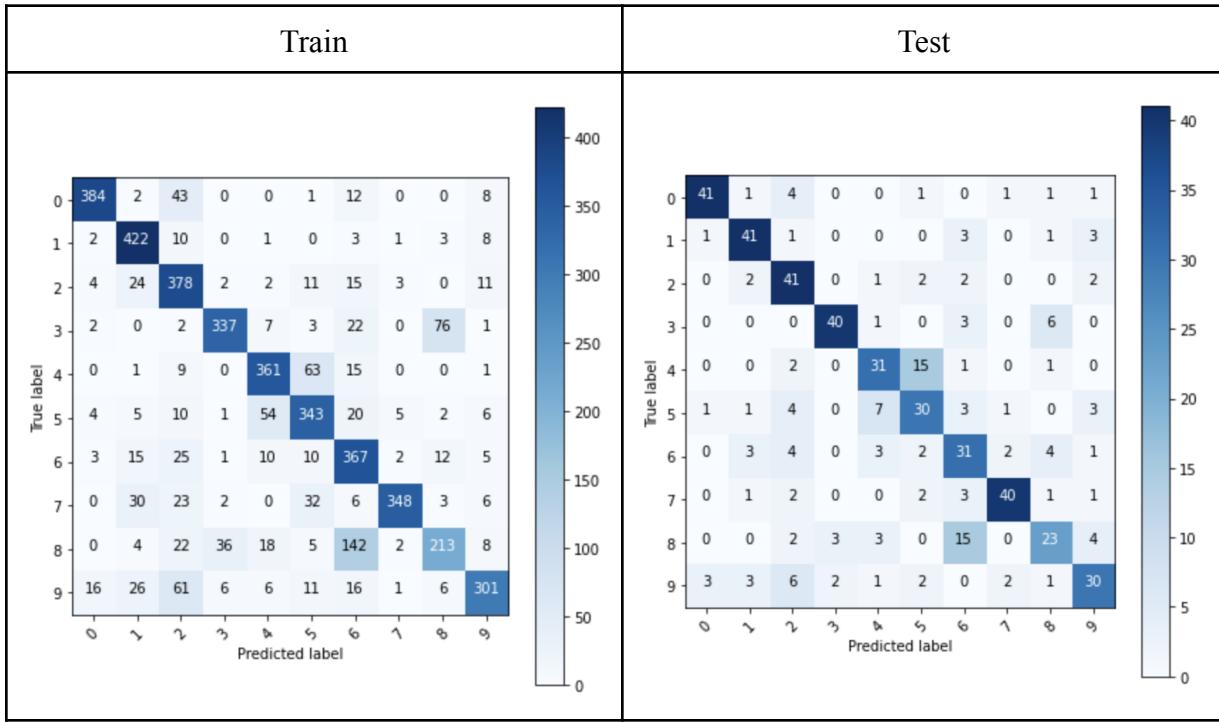
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	76%	69%
Loss	0.660881	0.876783

❖ Loss plot



## ❖ Confusion Matrix



### 1.3.2.2 Dropout = 0.7

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (9): ReLU(inplace=True)
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): ReLU(inplace=True)
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (19): ReLU(inplace=True)
```

```

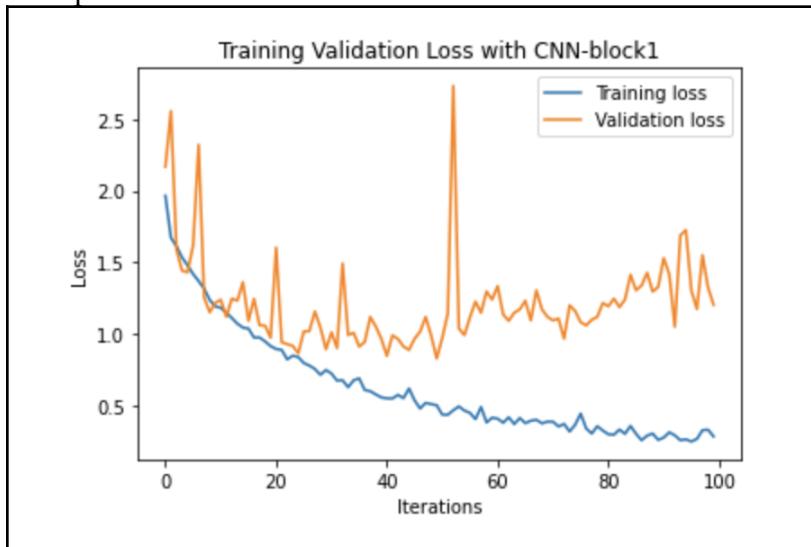
(20): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Dropout(p=0.7, inplace=False)
(1): Linear(in_features=64, out_features=10, bias=True)
)
)

```

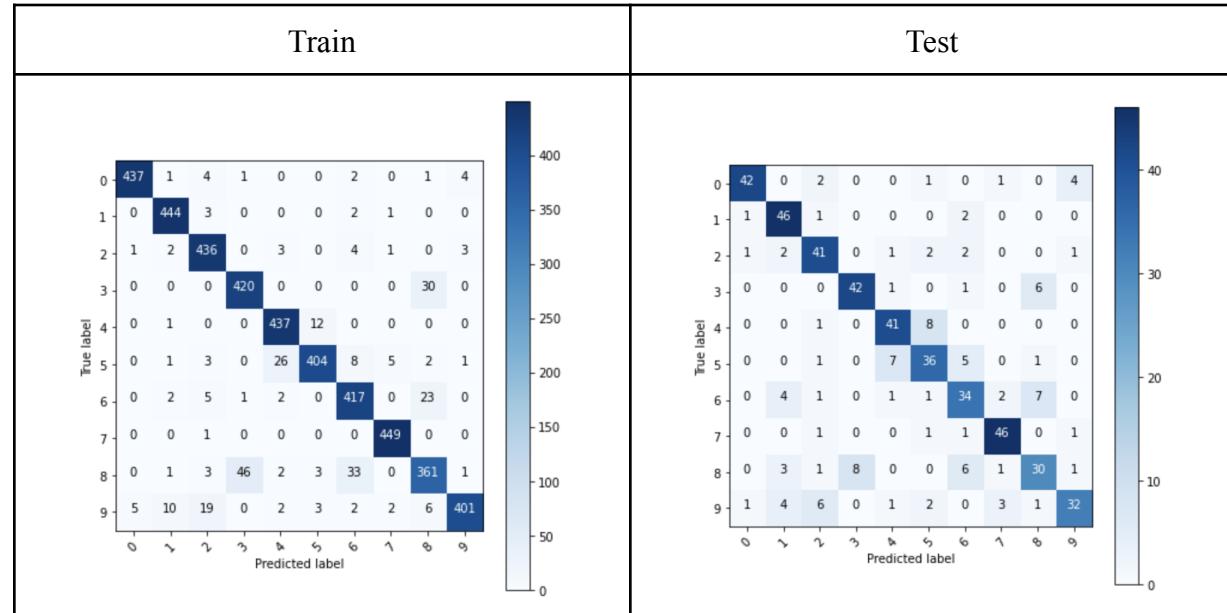
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	93%	78%
Loss	0.213234	0.833252

❖ Loss plot



❖ Confusion Matrix



### 1.3.3 Dropout after each block

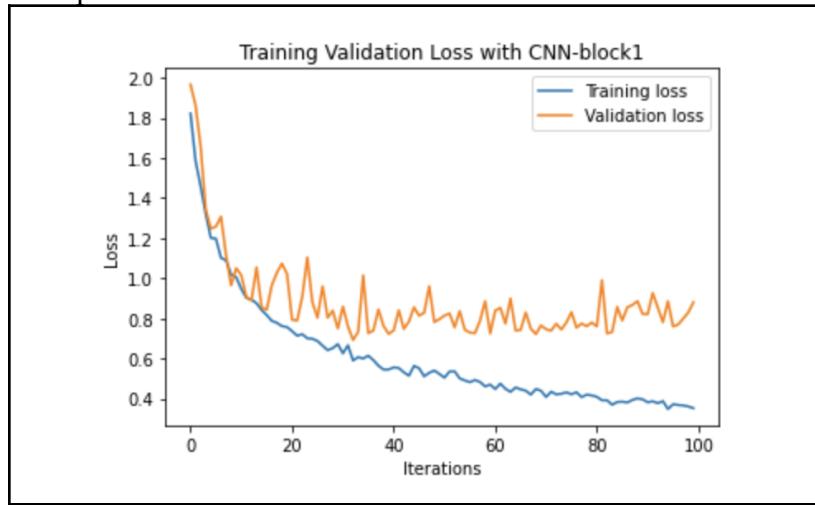
#### 1.3.3.1 Dropout = 0.3

```
Model_B123(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Dropout(p=0.3, inplace=False)
        (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (10): ReLU(inplace=True)
        (11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (13): ReLU(inplace=True)
        (14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): ReLU(inplace=True)
        (17): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (18): Dropout(p=0.3, inplace=False)
        (19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (21): ReLU(inplace=True)
        (22): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (23): Dropout(p=0.3, inplace=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)
```

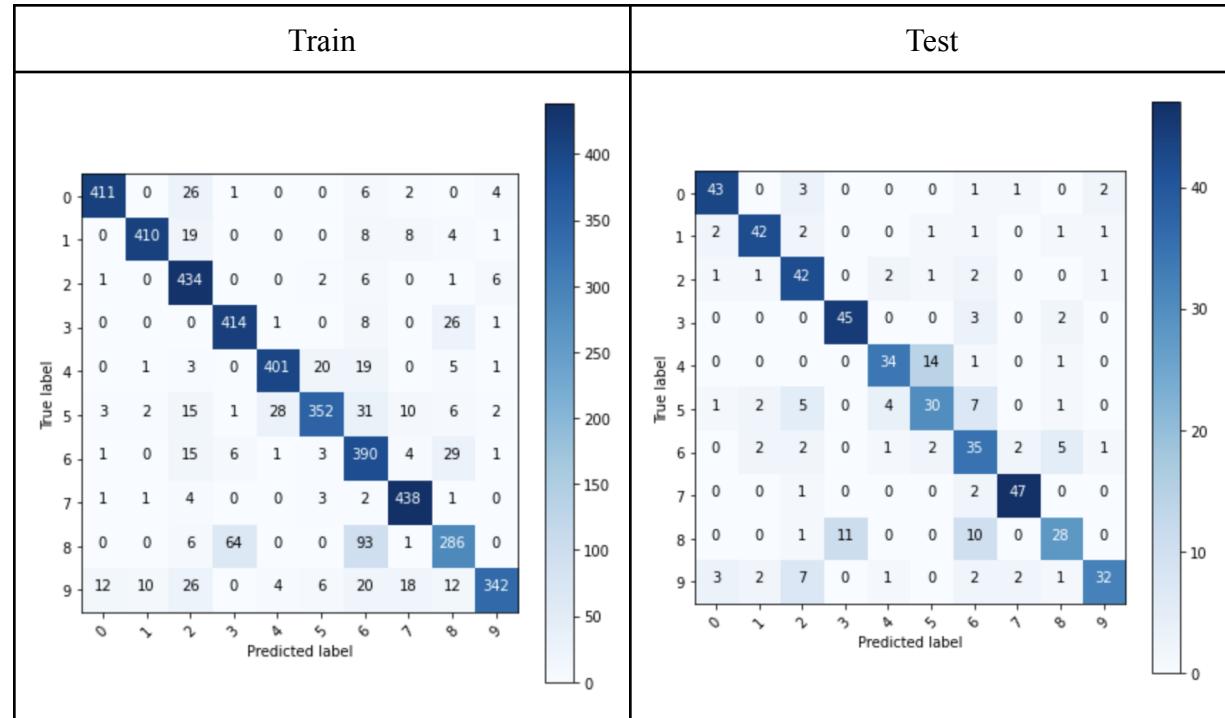
#### ❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	86%	75%
Loss	0.415009	0.690495

❖ Loss plot



❖ Confusion Matrix



### 1.3.3.2 Dropout = 0.5

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
```

```

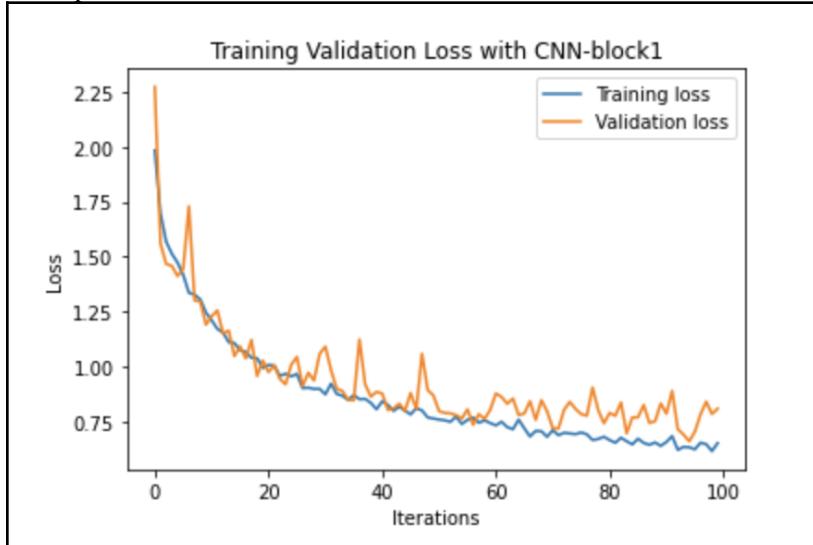
(7): Dropout(p=0.5, inplace=False)
(8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
(9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(10): ReLU(inplace=True)
(11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(13): ReLU(inplace=True)
(14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(16): ReLU(inplace=True)
(17): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
(18): Dropout(p=0.5, inplace=False)
(19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(21): ReLU(inplace=True)
(22): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
(23): Dropout(p=0.5, inplace=False)
)
(classifier): Sequential(
(0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

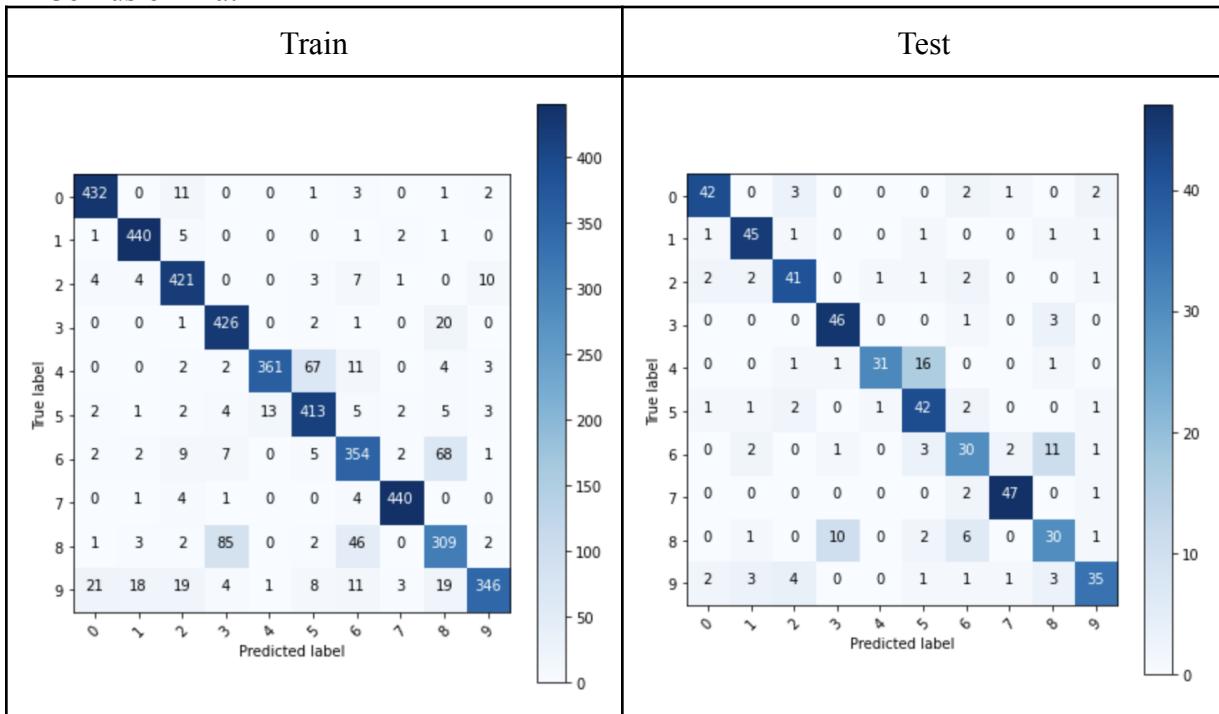
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	87%	77%
Loss	0.380031	0.662487

❖ Loss plot



## ❖ Confusion Matrix



### 1.3.3.3 Dropout = 0.7

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (7): Dropout(p=0.7, inplace=False)
        (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (10): ReLU(inplace=True)
        (11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (13): ReLU(inplace=True)
        (14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): ReLU(inplace=True)
        (17): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (18): Dropout(p=0.7, inplace=False)
        (19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
```

```

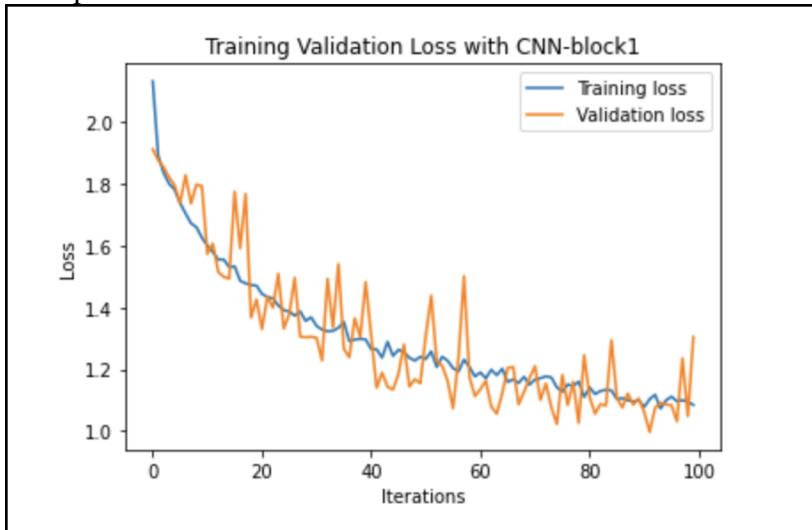
        (20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (21): ReLU(inplace=True)
        (22): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (23): Dropout(p=0.7, inplace=False)
    )
(classifier): Sequential(
    (0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

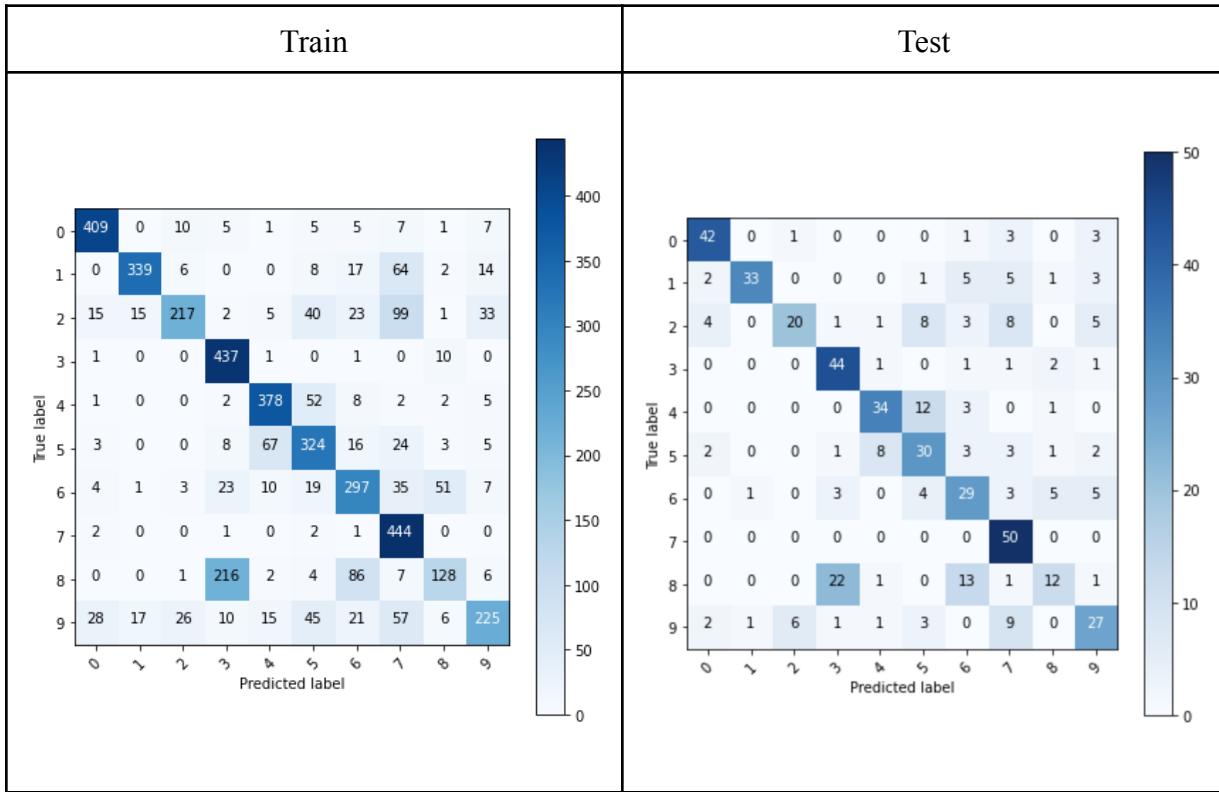
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	71%	64%
Loss	0.834854	0.997905

❖ Loss plot



## ❖ Confusion Matrix



### Analysis:

- Dropout was implemented in all three cases with 3 dropout probabilities in each case.
- Inclusion of dropout prevented overfitting. The extent of the prevention depended on the dropout probability.
- In every case, 0.3 dropout probability still tends to overfit while 0.7 dropout probability tends to underfit. Therefore, 0.5 dropout probability is the best out of the 3 probabilities tested.
- Out of the 3 architectures, best results are obtained if dropout is applied after each block.
- Dropout applied after convolutional layers is not successful, as we might lose important information in images by dropping out some nodes also CNN already has shared weights so dropping out can have adverse effects.
- From the graphs it is clearly visible at 0.3 after some point test loss keeps on increasing while train loss keeps on increasing i.e overfitting. In our case 0.5 works the best with between the blocks.

## 1.4 Hyperparameters and best architecture

### 1.4.1 Number of layers (X,Y,Z)

#### 1.4.1.1. X = 2, Y = 2, Z = 2

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
```

```

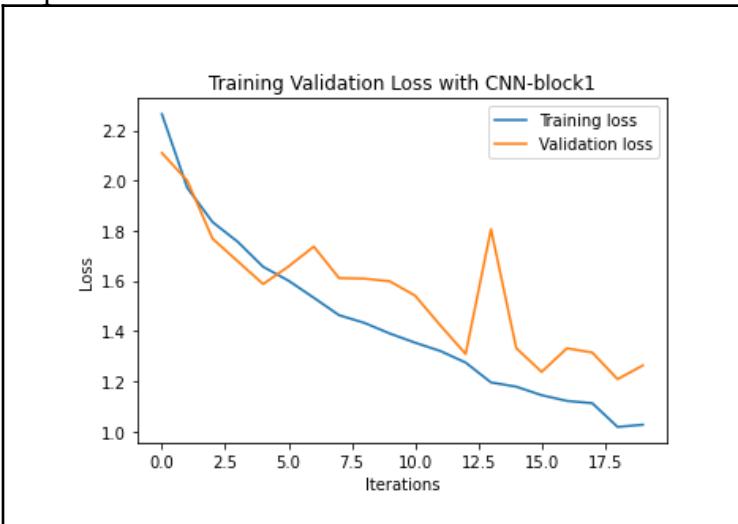
track_running_stats=True)
(5): ReLU(inplace=True)
(6): MaxPool2d(kernel_size=2, stride=1, padding=0, dilation=1,
ceil_mode=False)
(7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
(8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(9): ReLU(inplace=True)
(10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(12): ReLU(inplace=True)
(13): MaxPool2d(kernel_size=2, stride=1, padding=0,
dilation=1, ceil_mode=False)
(14): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(15): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(16): ReLU(inplace=True)
(17): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1))
(18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(19): ReLU(inplace=True)
(20): MaxPool2d(kernel_size=4, stride=3, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Linear(in_features=1600, out_features=512, bias=True)
(1): ReLU(inplace=True)
(2): Linear(in_features=512, out_features=256, bias=True)
(3): ReLU(inplace=True)
(4): Linear(in_features=256, out_features=64, bias=True)
(5): ReLU(inplace=True)
(6): Linear(in_features=64, out_features=10, bias=True)
)
)

```

#### ❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	65%	55%
Loss	1.011410	1.203690

❖ Loss plot



1.4.1.2. X = 3, Y = 2, Z = 1

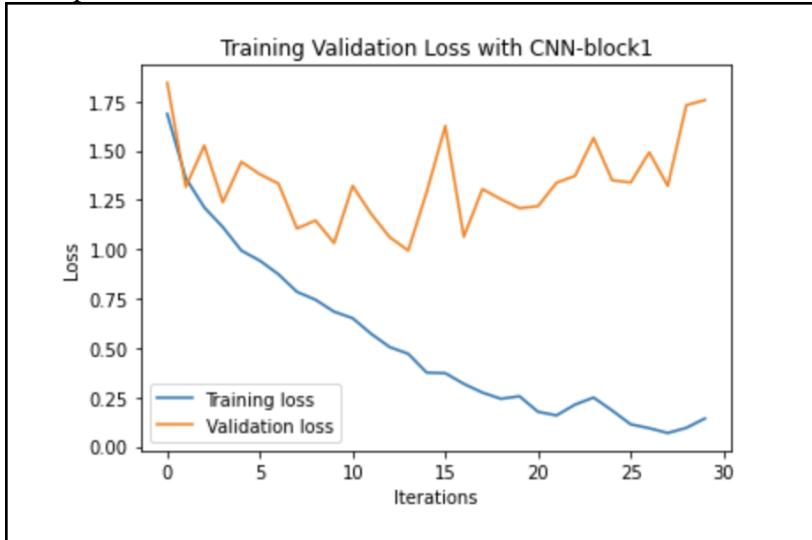
```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (7): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (8): ReLU(inplace=True)
        (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
        (10): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): ReLU(inplace=True)
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): ReLU(inplace=True)
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
        (17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (19): ReLU(inplace=True)
        (20): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)
```

)

❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	87%	66%
Loss	0.376737	0.995357

❖ Loss plot



#### 1.4.1.3. X = 2, Y = 3, Z = 1

```
Model_B123(  
    features): Sequential(  
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))  
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (2): ReLU(inplace=True)  
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))  
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (5): ReLU(inplace=True)  
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
            ceil_mode=False)  
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))  
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (9): ReLU(inplace=True)  
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))  
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (12): ReLU(inplace=True)  
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))  
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)
```

```

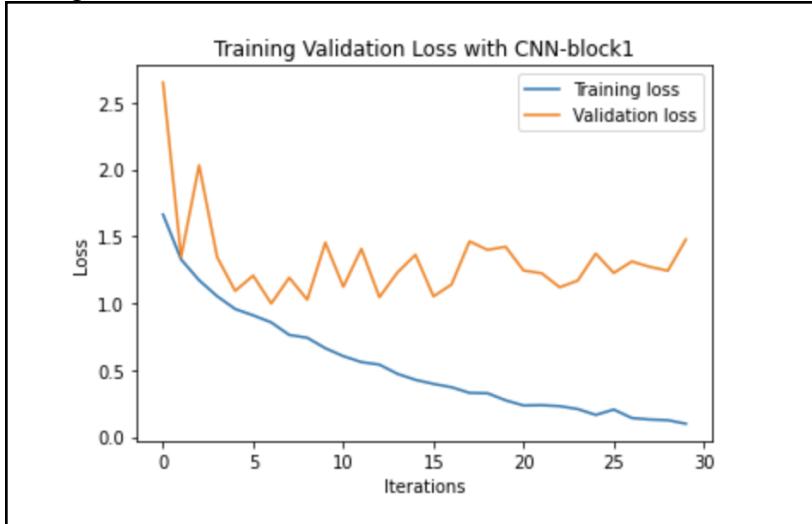
(15): ReLU(inplace=True)
(16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
(17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(19): ReLU(inplace=True)
(20): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	74%	64%
Loss	0.718399	0.995629

❖ Loss plot



Best value: X = 2, Y = 3, Z = 1

- The architecture gave best testing accuracy and minimum loss.
- X = 2, Y = 3, Z = 1 and X = 3, Y = 2, Z = 1 are almost similar so we chose first one.
- Keeping higher X value worked better as it helped in better feature extraction. Hence, X = 1 gave low accuracy while X = 2 and X = 3 gave better accuracy.

#### 1.4.2 Batch Normalization (With Max Pooling)

##### 1.4.2.1. Without Batch Normalization

```

Model_B123(
(features): Sequential(

```

```

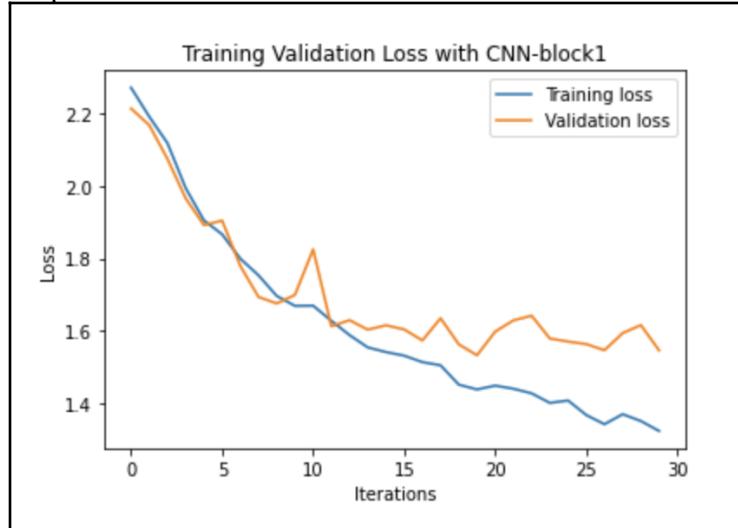
(0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
(1): ReLU(inplace=True)
(2): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
(3): ReLU(inplace=True)
(4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(5): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
(6): ReLU(inplace=True)
(7): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(8): ReLU(inplace=True)
(9): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(10): ReLU(inplace=True)
(11): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
(12): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(13): ReLU(inplace=True)
(14): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	52%	45%
Loss	1.349049	1.532184

❖ Loss plot



#### 1.4.2.2. With Batch Normalization

```

Model_B123(
(features): Sequential(

```

```

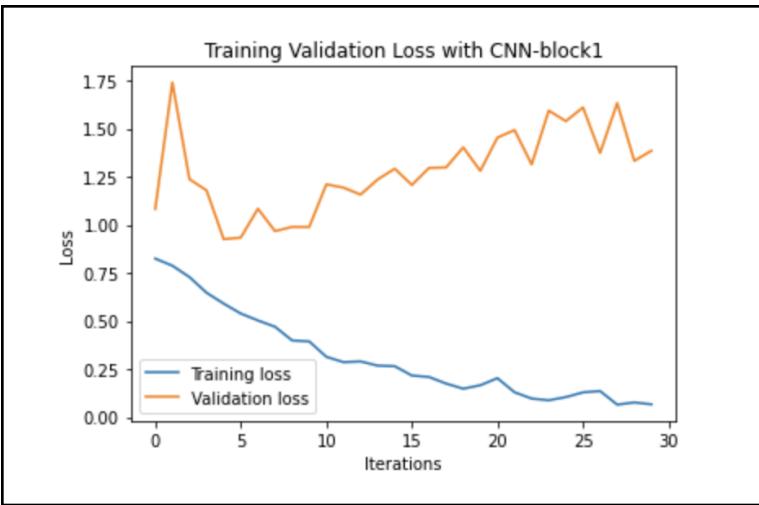
(0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
(1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(2): ReLU(inplace=True)
(3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
(4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(5): ReLU(inplace=True)
(6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
(7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
(8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(9): ReLU(inplace=True)
(10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(12): ReLU(inplace=True)
(13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(15): ReLU(inplace=True)
(16): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
(17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(19): ReLU(inplace=True)
(20): MaxPool2d(kernel_size=2, stride=2, padding=0,
dilation=1, ceil_mode=False)
)
(classifier): Sequential(
(0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	81%	69%
Loss	0.521781	0.919671

❖ Loss plot



**Best value:** The model performed better with batch normalization as it normalizes value in a range (0,1) which helps the model learn faster and perform better.

### 1.4.3 Pooling

#### 1.4.3.1. Max Pooling

The model is the same as that in section 1.4.2.2 (with batch normalization and max pooling).

#### 1.4.3.2. Average Pooling

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (7): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (8): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (9): ReLU(inplace=True)
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): ReLU(inplace=True)
        (13): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (14): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (15): ReLU(inplace=True)
        (16): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (17): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (18): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (19): ReLU(inplace=True)
        (20): AvgPool2d(kernel_size=2, stride=2, padding=0)
```

```

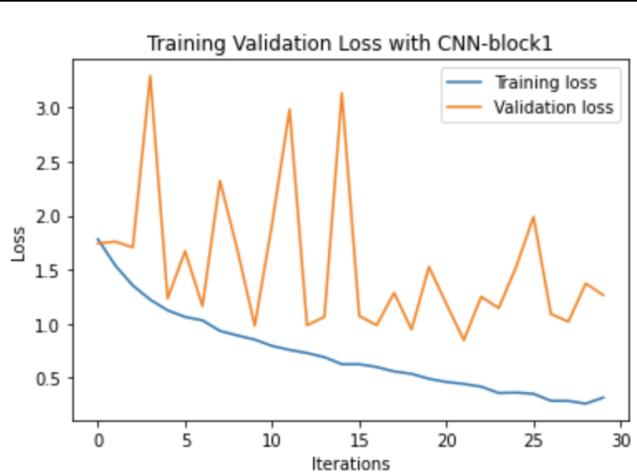
)
(classifier): Sequential(
    (0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	88%	72%
Loss	0.347223	0.853640

❖ Loss plot



**Best value:** The model performed better with Average Pooling because Average Pooling considers every feature of the image. Looking at the dataset images, the sharpness and intensity in different areas of images is different. Therefore, average pooling takes into account all of them rather than the highest (like Max Pooling does).

#### 1.4.4 Dropout

From the analysis of section 1.3, the optimal dropout probability was chosen as 0.5 and the dropout was applied after each block for best performance.

```

Model_B123(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (7): Dropout(p=0.5, inplace=False)
        (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
    )
)

```

```

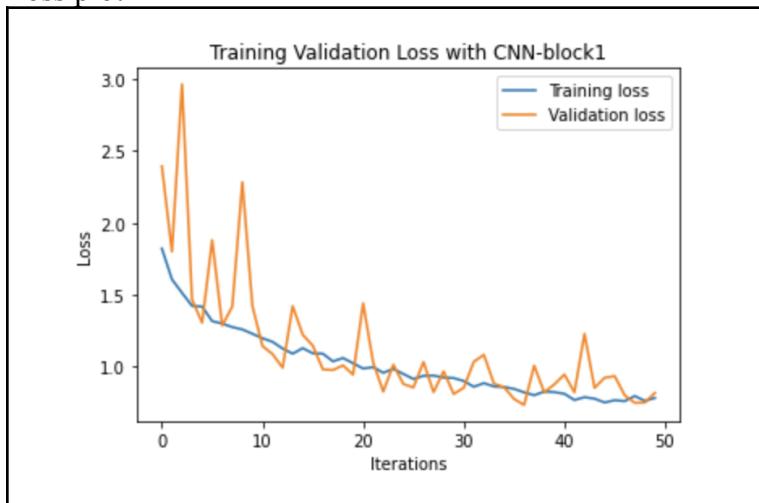
(9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(10): ReLU(inplace=True)
(11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(13): ReLU(inplace=True)
(14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(16): ReLU(inplace=True)
(17): AvgPool2d(kernel_size=2, stride=2, padding=0)
(18): Dropout(p=0.5, inplace=False)
(19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(21): ReLU(inplace=True)
(22): AvgPool2d(kernel_size=2, stride=2, padding=0)
(23): Dropout(p=0.5, inplace=False)
)
(classifier): Sequential(
(0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

#### ❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	81%	73%
Loss	0.529796	0.738360

#### ❖ Loss plot



#### 1.4.5 Optimizer and weight decay

Adam optimizer has been chosen for the model as Adam generally is the optimal choice because of its adaptive learning rate property.

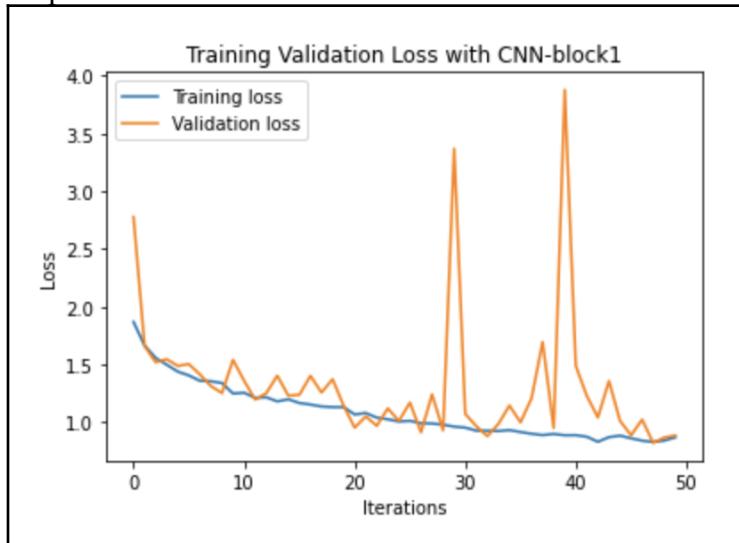
#### 1.4.5.1 With weight decay

```
Model_B123(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (7): Dropout(p=0.5, inplace=False)
        (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (10): ReLU(inplace=True)
        (11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (13): ReLU(inplace=True)
        (14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): ReLU(inplace=True)
        (17): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (18): Dropout(p=0.5, inplace=False)
        (19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (21): ReLU(inplace=True)
        (22): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (23): Dropout(p=0.5, inplace=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)
```

#### ❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	78%	71%
Loss	0.598614	0.823383

❖ Loss plot



#### 1.4.5.2. Without Weight Decay

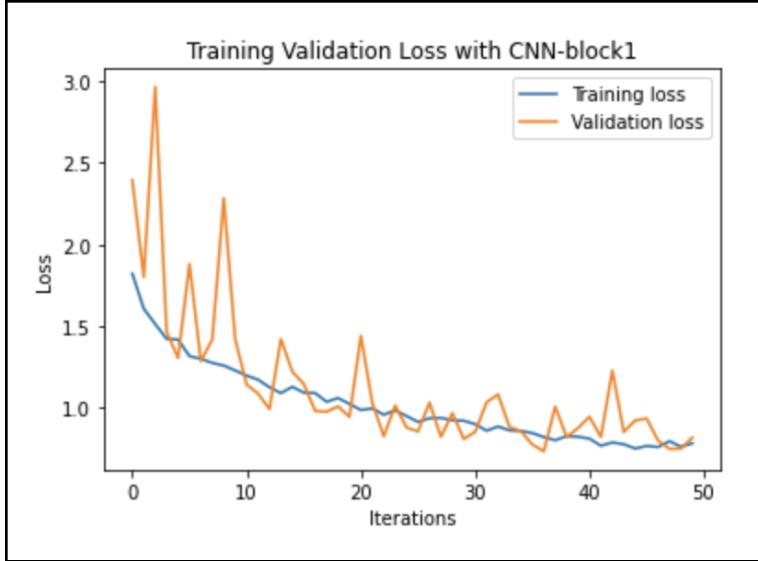
```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (7): Dropout(p=0.5, inplace=False)
        (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (10): ReLU(inplace=True)
        (11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (13): ReLU(inplace=True)
        (14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): ReLU(inplace=True)
        (17): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (18): Dropout(p=0.5, inplace=False)
        (19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (21): ReLU(inplace=True)
        (22): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (23): Dropout(p=0.5, inplace=False)
    )
    classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
```

```
)  
)
```

❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	81%	73%
Loss	0.529796	0.738360

❖ Loss plot



**Best value:** Adam optimizer performed well. Using weight decay did not improve performance. If the value of weight decay was kept small, no improvement was seen whereas if weight decay was kept large then the model underfit. Therefore, no weight decay was done.

## 1.4.6 Learning Rate

### 1.4.6.1 Learning rate = 0.01

The results are the same as section 1.4.5.2 as the learning rate was 0.01 by default in all of the above results.

### 1.4.6.2. Learning rate = 0.001

```
Model_B123(  
    features): Sequential(  
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))  
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (2): ReLU(inplace=True)  
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))  
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,  
            track_running_stats=True)  
        (5): ReLU(inplace=True)  
        (6): AvgPool2d(kernel_size=2, stride=2, padding=0)
```

```

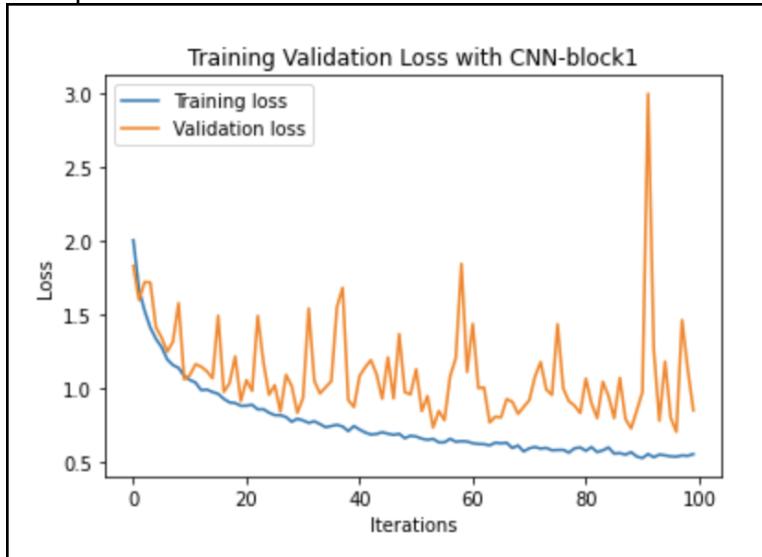
(7): Dropout(p=0.5, inplace=False)
(8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
(9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(10): ReLU(inplace=True)
(11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(13): ReLU(inplace=True)
(14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
(15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(16): ReLU(inplace=True)
(17): AvgPool2d(kernel_size=2, stride=2, padding=0)
(18): Dropout(p=0.5, inplace=False)
(19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(21): ReLU(inplace=True)
(22): AvgPool2d(kernel_size=2, stride=2, padding=0)
(23): Dropout(p=0.5, inplace=False)
)
(classifier): Sequential(
(0): Linear(in_features=64, out_features=10, bias=True)
)
)

```

#### ❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	86%	76%
Loss	0.370939	0.703236

#### ❖ Loss plot



**Best value:** Learning rate = 0.001 performed better as, at value 0.01, the model oscillated around the minima and could not reach it which indicated that the value had to be decreased.

## 1.4.7 Data Augmentation

### 1.4.7.1 Without Data Augmentation

The results are the same as section 1.4.6.2 as the model in the section did not use data augmentation and had optimal hyperparameters.

### 1.4.7.2. With Data Augmentation

Random Horizontal flips were used for data augmentation.

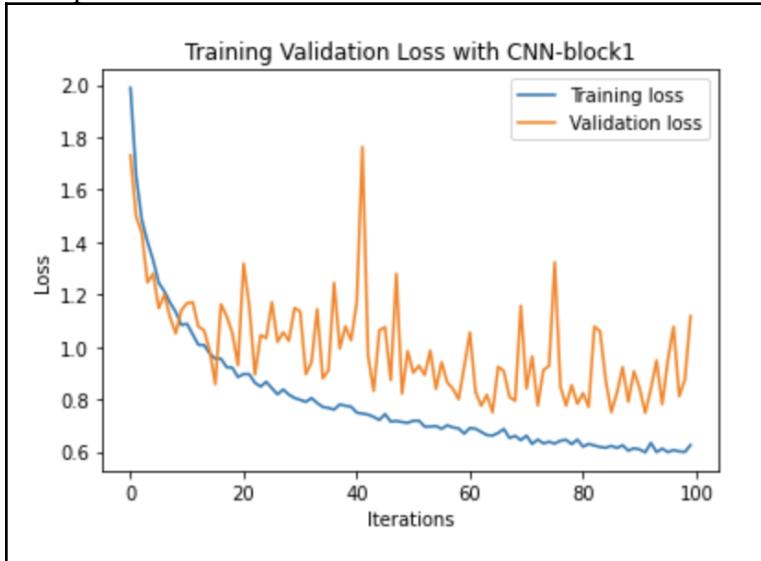
```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (7): Dropout(p=0.5, inplace=False)
        (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (10): ReLU(inplace=True)
        (11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (13): ReLU(inplace=True)
        (14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): ReLU(inplace=True)
        (17): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (18): Dropout(p=0.5, inplace=False)
        (19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (21): ReLU(inplace=True)
        (22): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (23): Dropout(p=0.5, inplace=False)
    )
    classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)
```

#### ❖ Accuracy and Loss

Block 1,2,3	Train	Test
-------------	-------	------

Accuracy	83%	74%
Loss	0.461166	0.749098

❖ Loss plot



**Best value:** Data augmentation improved the performance of the model by making it more generalised. It made the model more generic and helped the model make more accurate predictions.

#### 1.4.8 Epochs

The model was run for 500 epochs.

```
Model_B123(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (7): Dropout(p=0.5, inplace=False)
        (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (10): ReLU(inplace=True)
        (11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (13): ReLU(inplace=True)
        (14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): ReLU(inplace=True)
    )
)
```

```

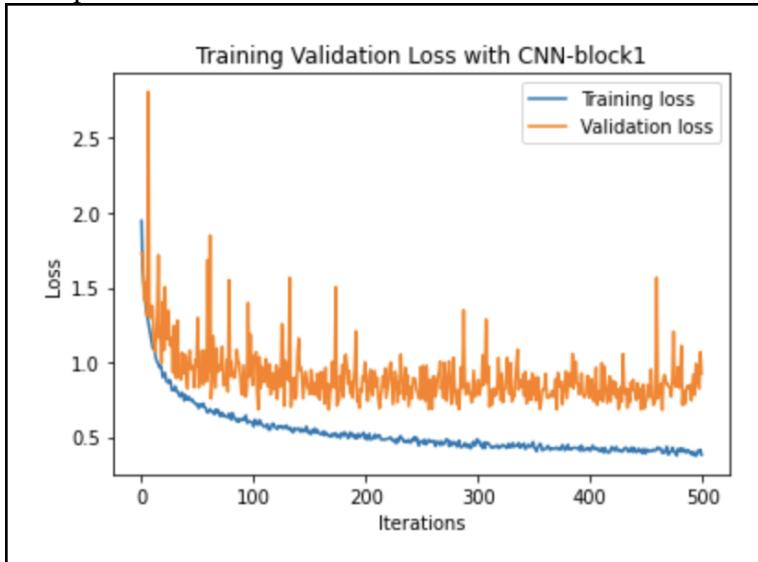
(17) : AvgPool2d(kernel_size=2, stride=2, padding=0)
(18) : Dropout(p=0.5, inplace=False)
(19) : Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
(20) : BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
(21) : ReLU(inplace=True)
(22) : AvgPool2d(kernel_size=2, stride=2, padding=0)
(23) : Dropout(p=0.5, inplace=False)
)
(classifier) : Sequential(
(0) : Linear(in_features=64, out_features=10, bias=True)
)
)
)

```

❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	92%	78%
Loss	0.207575	0.683457

❖ Loss plot



**Best value:** From the plot and model performance, it was found that optimal epoch value was 100. After that, the model learning saturates i.e there is no significant change in test loss.

**Best model:**

Layers: X = 2, Y = 3, Z = 1

Batch Normalization

Average Pooling

Dropout = 0.5 after each block

Adam optimizer

No weight decay

Learning rate = 0.001

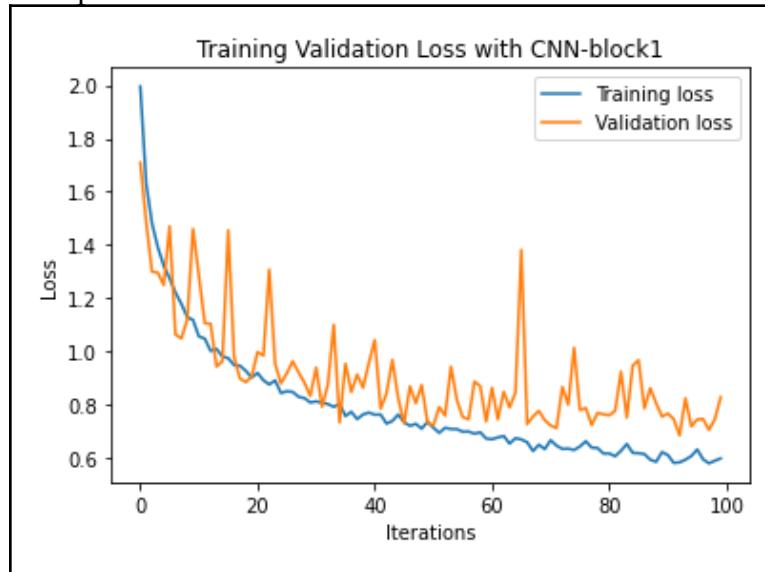
Data Augmentation - Random Horizontal Flips  
Epochs = 100

```
Model_B123(
    (features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (4): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (7): Dropout(p=0.5, inplace=False)
        (8): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (10): ReLU(inplace=True)
        (11): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (12): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (13): ReLU(inplace=True)
        (14): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (15): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): ReLU(inplace=True)
        (17): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (18): Dropout(p=0.5, inplace=False)
        (19): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (20): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (21): ReLU(inplace=True)
        (22): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (23): Dropout(p=0.5, inplace=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)
```

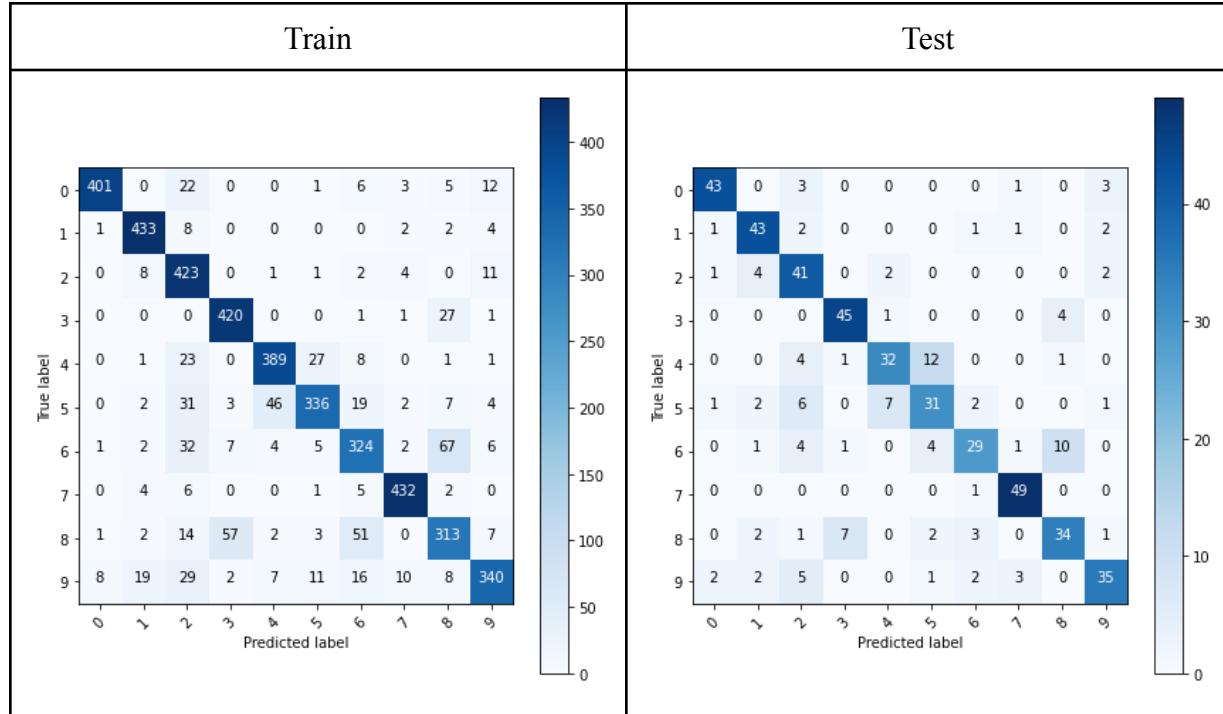
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	84%	76%
Loss	0.420687	0.682963

❖ Loss plot



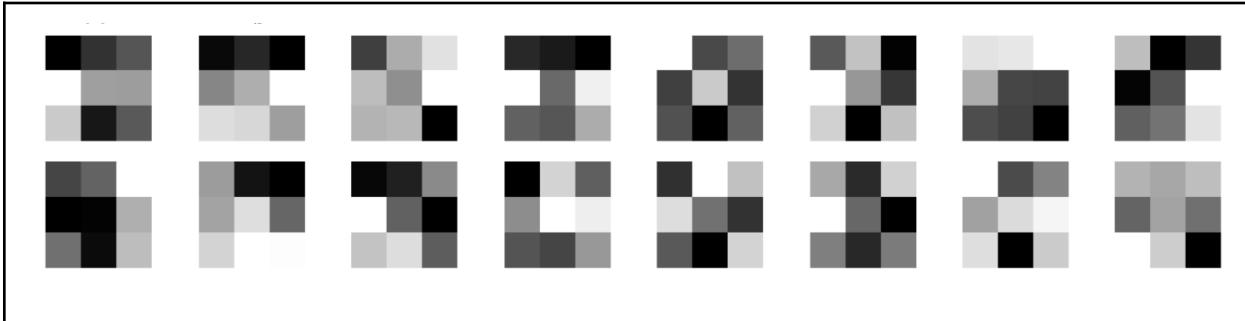
❖ Confusion Matrix



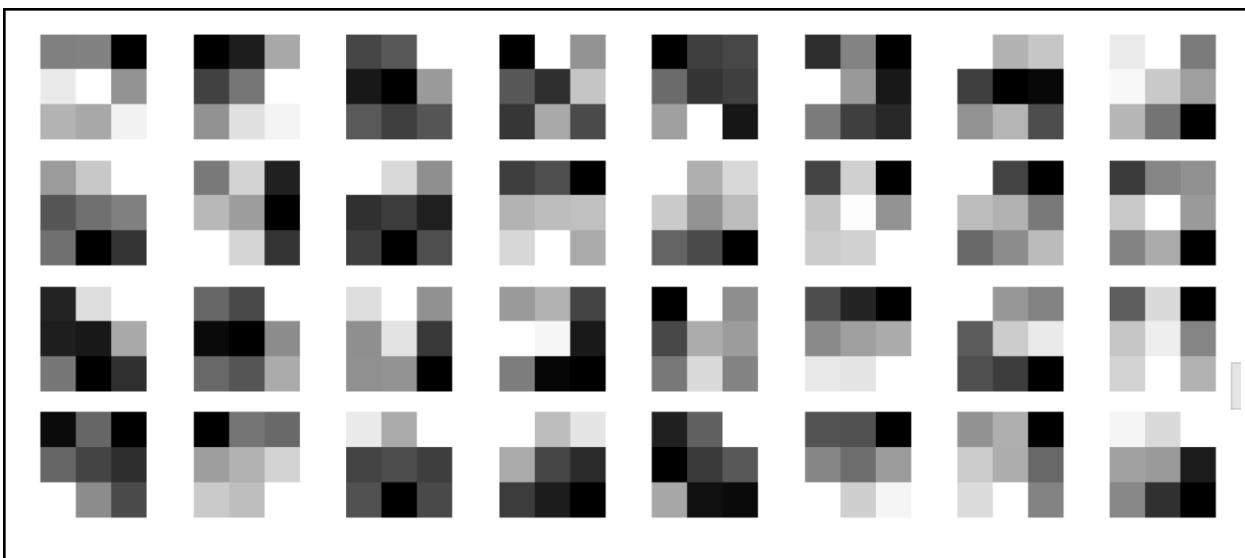
## 1.5 Visualize Convolutional Filters and Feature Maps

### 1.5.1 Visualize Convolutional Filters

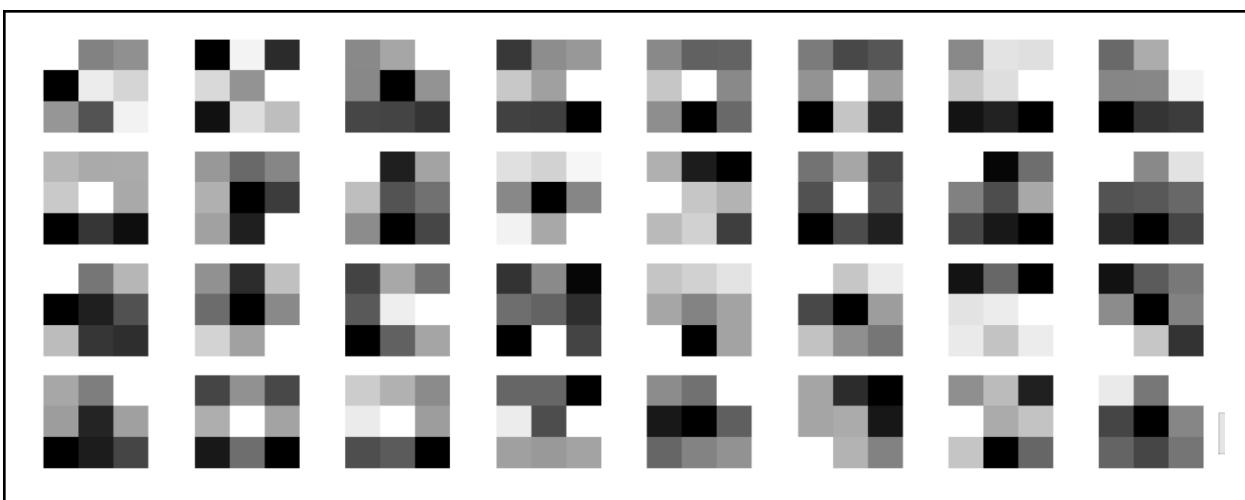
- Block 1 filter:



- Block 2 filter:

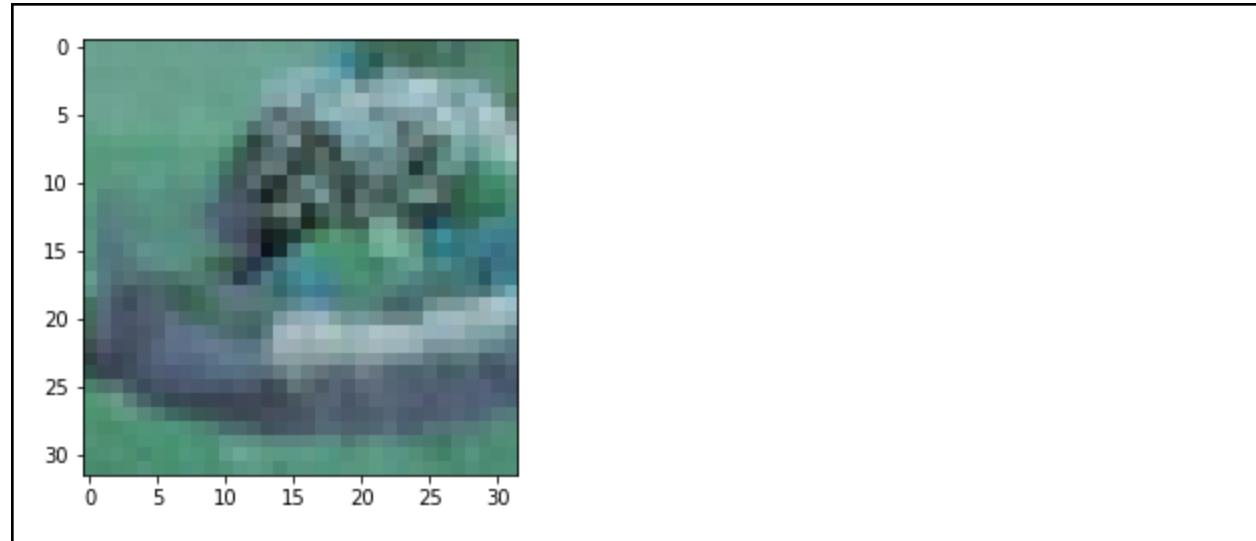


- Block 3 filter:

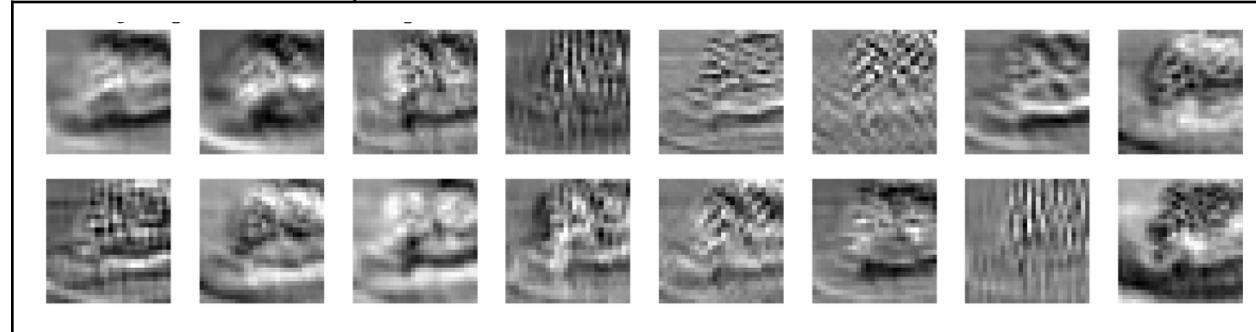


### 1.5.2 Visualize Feature Maps

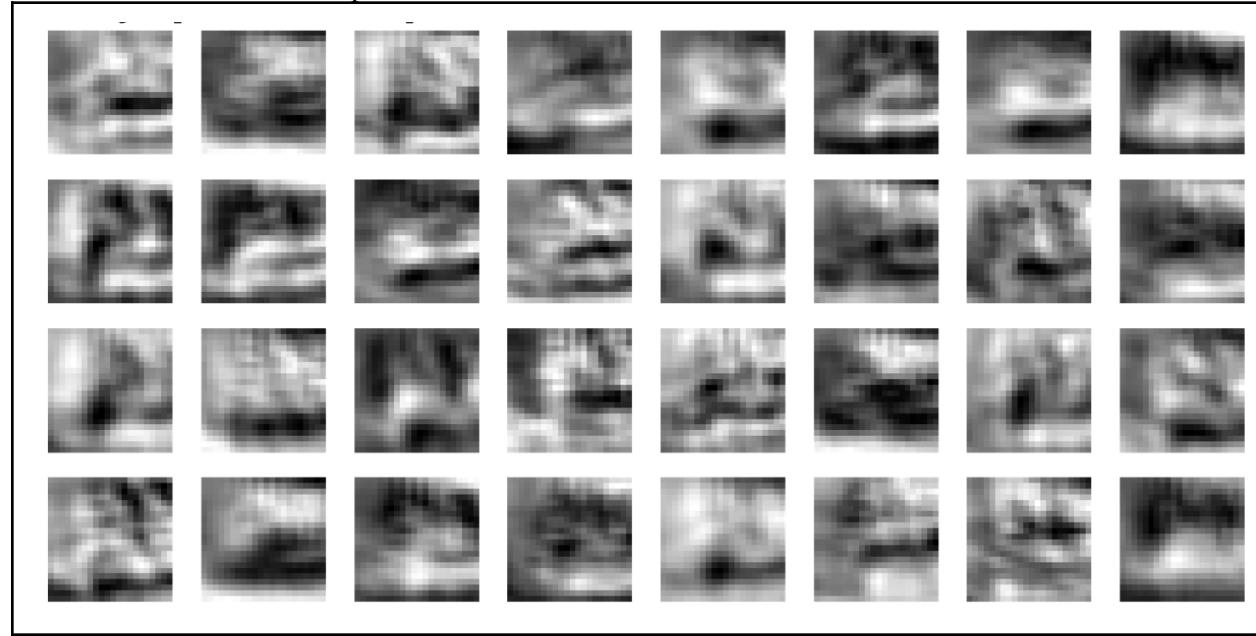
For the following image feature maps visualised are:



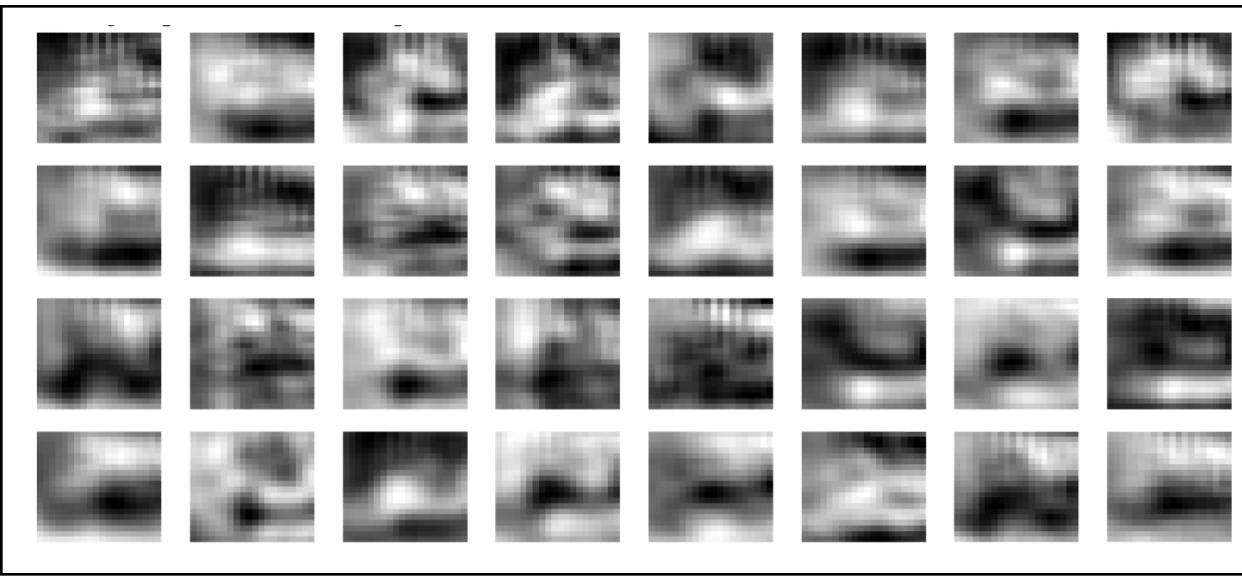
- Block 1 feature maps:



- Block 2 feature maps:



- Some of the block 3 feature maps:



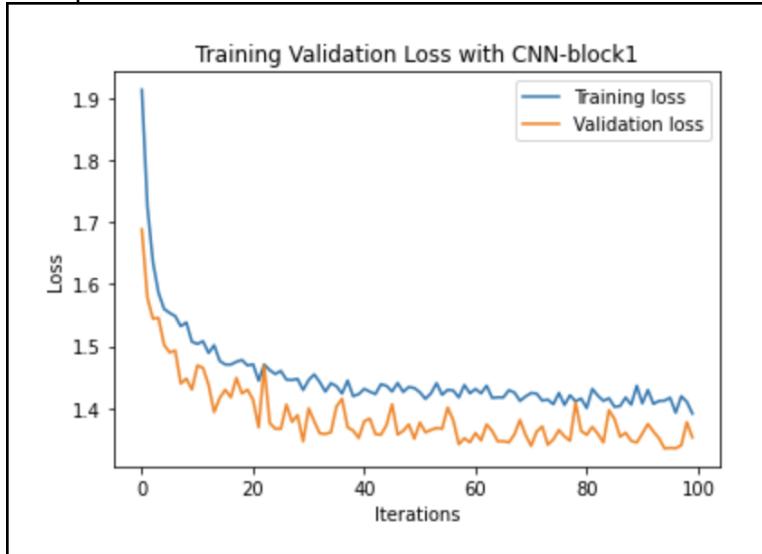
## 1.6 Best Model Without Activation Function

```
Model_B123(
    features): Sequential(
        (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1))
        (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (2): Conv2d(16, 16, kernel_size=(3, 3), stride=(1, 1))
        (3): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (4): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (5): Dropout(p=0.5, inplace=False)
        (6): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1))
        (7): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (8): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (9): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (10): Conv2d(32, 32, kernel_size=(3, 3), stride=(1, 1))
        (11): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (12): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (13): Dropout(p=0.5, inplace=False)
        (14): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1))
        (15): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (16): AvgPool2d(kernel_size=2, stride=2, padding=0)
        (17): Dropout(p=0.5, inplace=False)
    )
    (classifier): Sequential(
        (0): Linear(in_features=64, out_features=10, bias=True)
    )
)
```

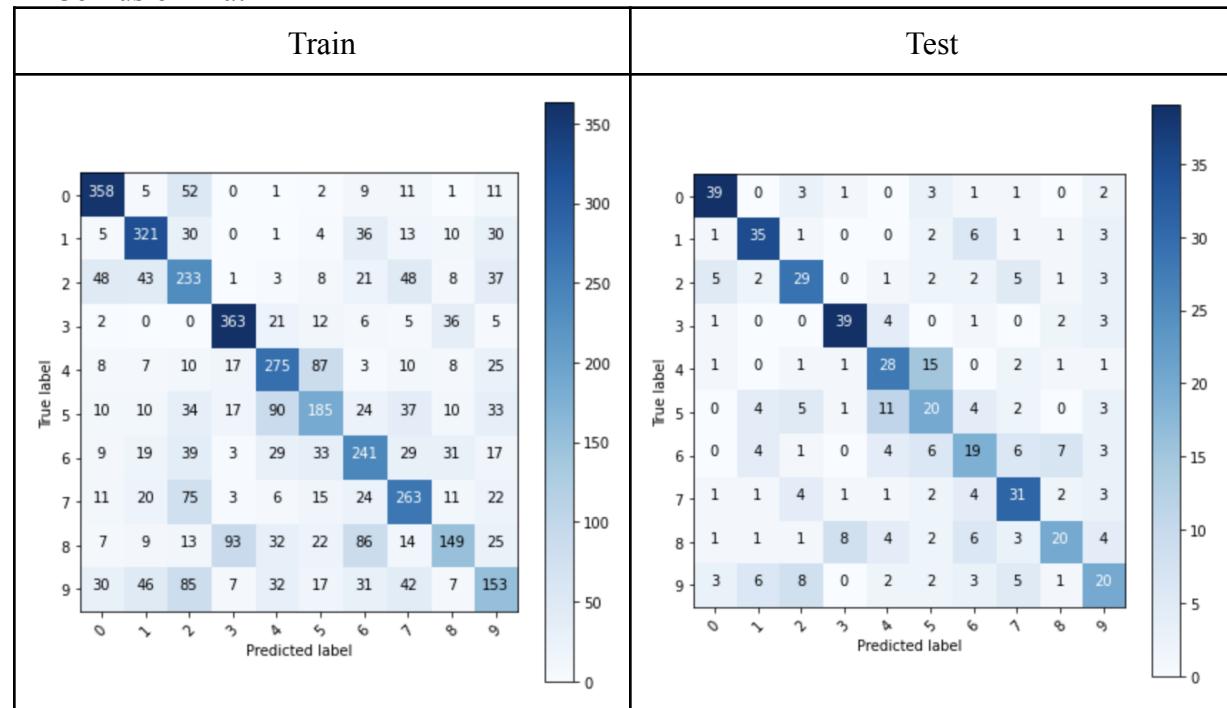
❖ Accuracy and Loss

Block 1,2,3	Train	Test
Accuracy	56%	56%
Loss	1.278515	1.338608

❖ Loss plot



❖ Confusion Matrix



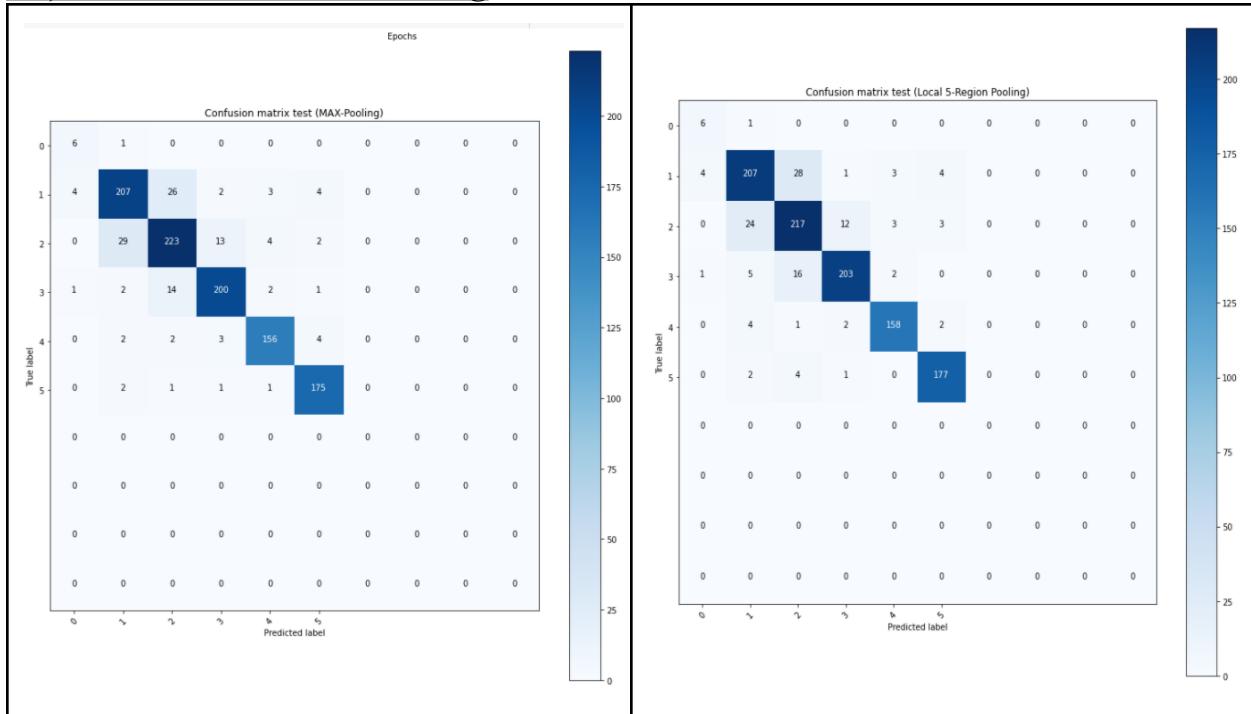
**Analysis:**

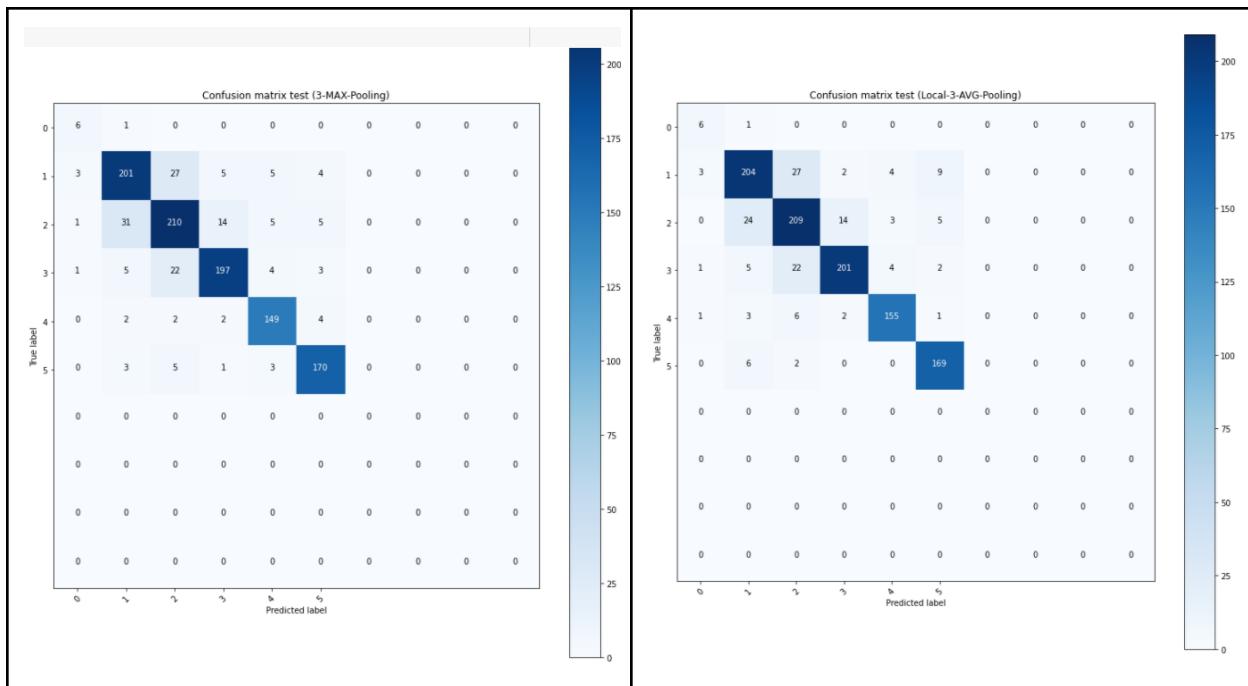
- Removal of activation function has a significant drop in accuracy (from 76% to 56%) on test data.
- The drop in performance is because, without activation function, the model cannot capture non-linearity in the data.
- The model also might have a vanishing/exploding gradient problem without activation function.

## 2. Sentence Classification

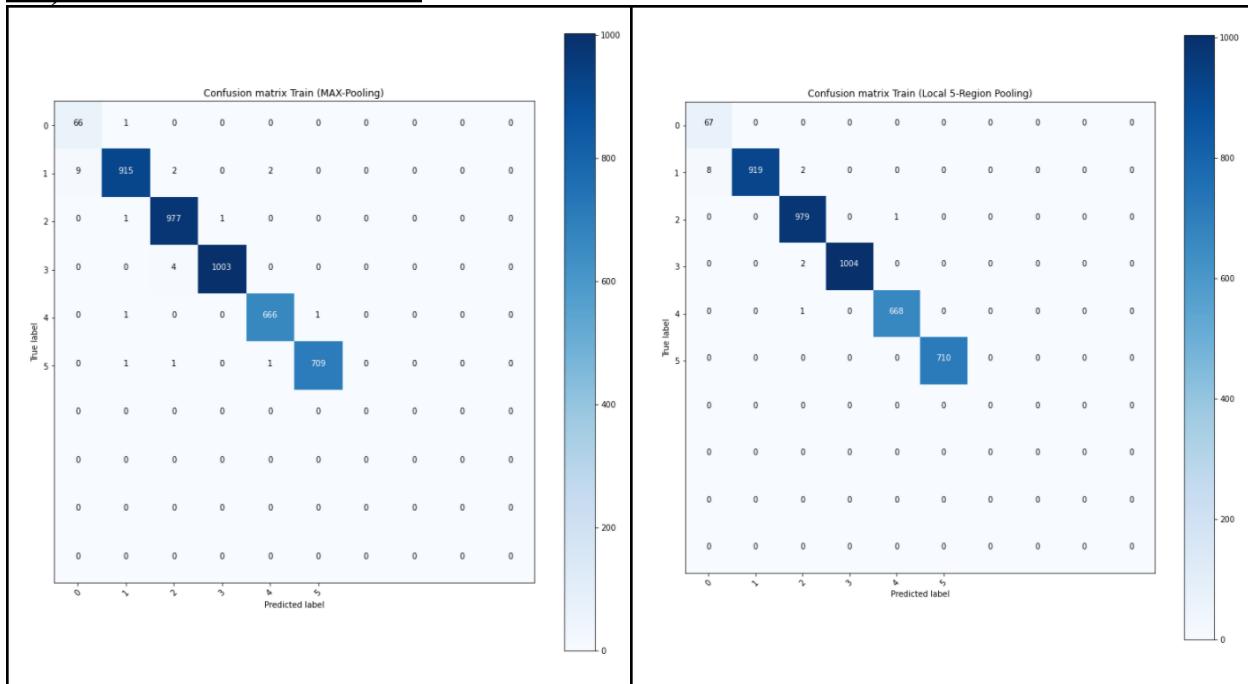
Description	Values
input word vectors	Google word2vec
filter region size	(3,4,5)
feature maps	100
activation function	ReLU
pooling	1-max pooling
dropout rate	0.5
$l_2$ norm constraint	3

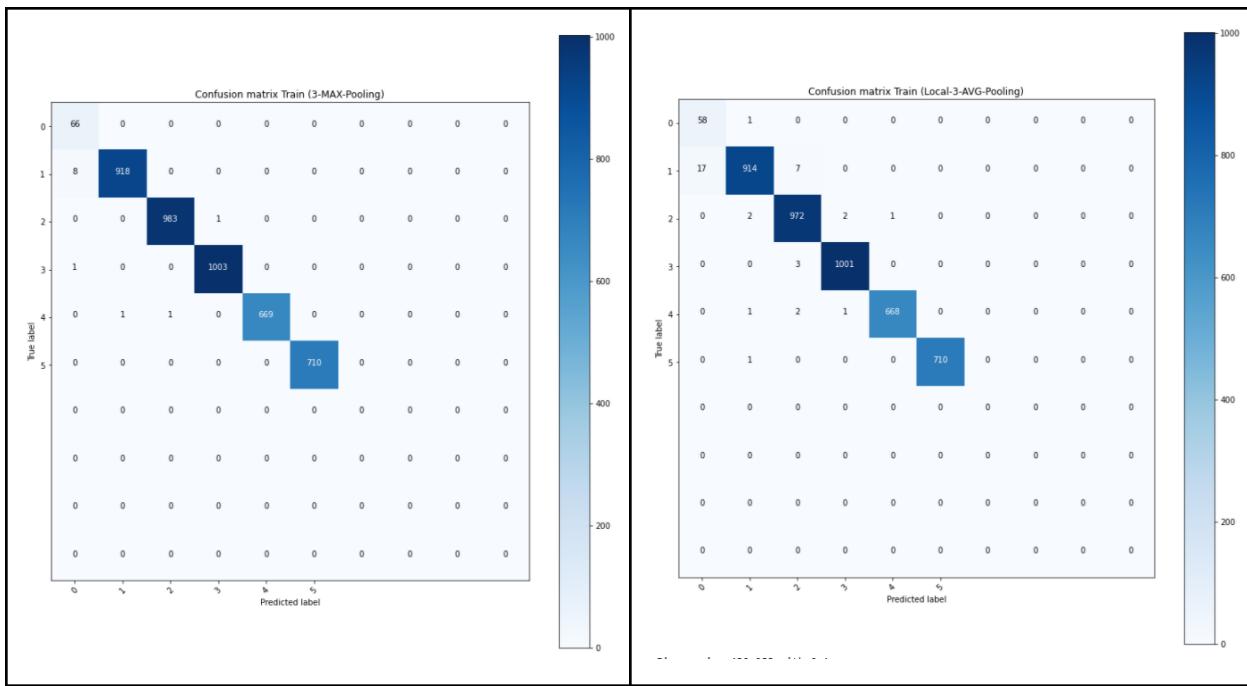
### a.1) Confusion matrix - Training



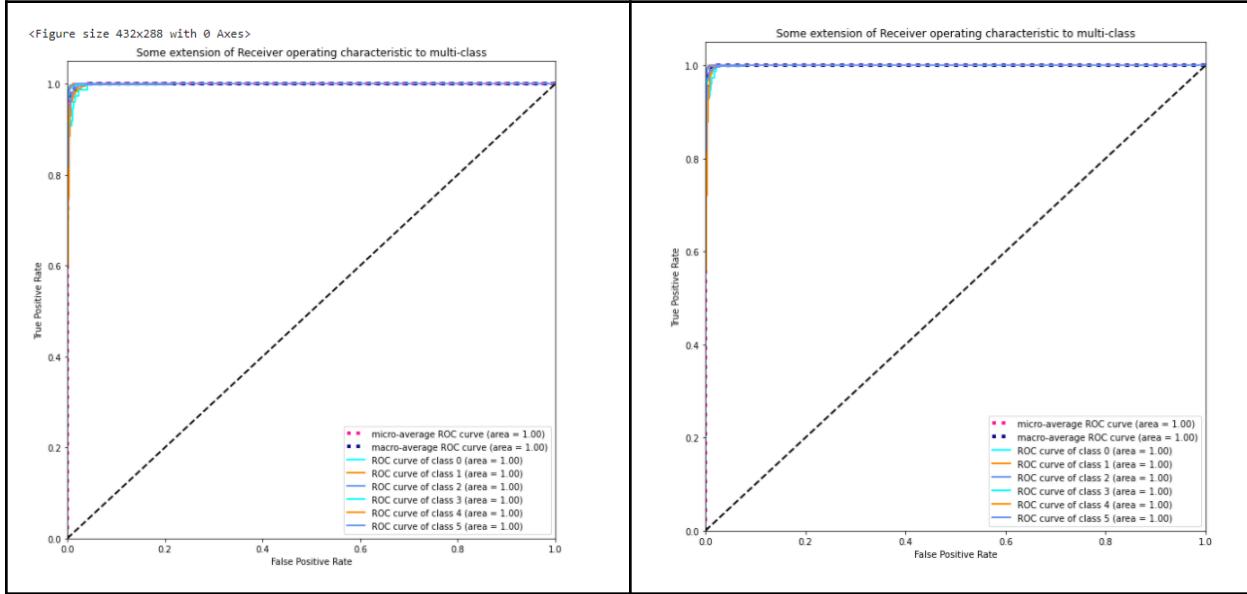


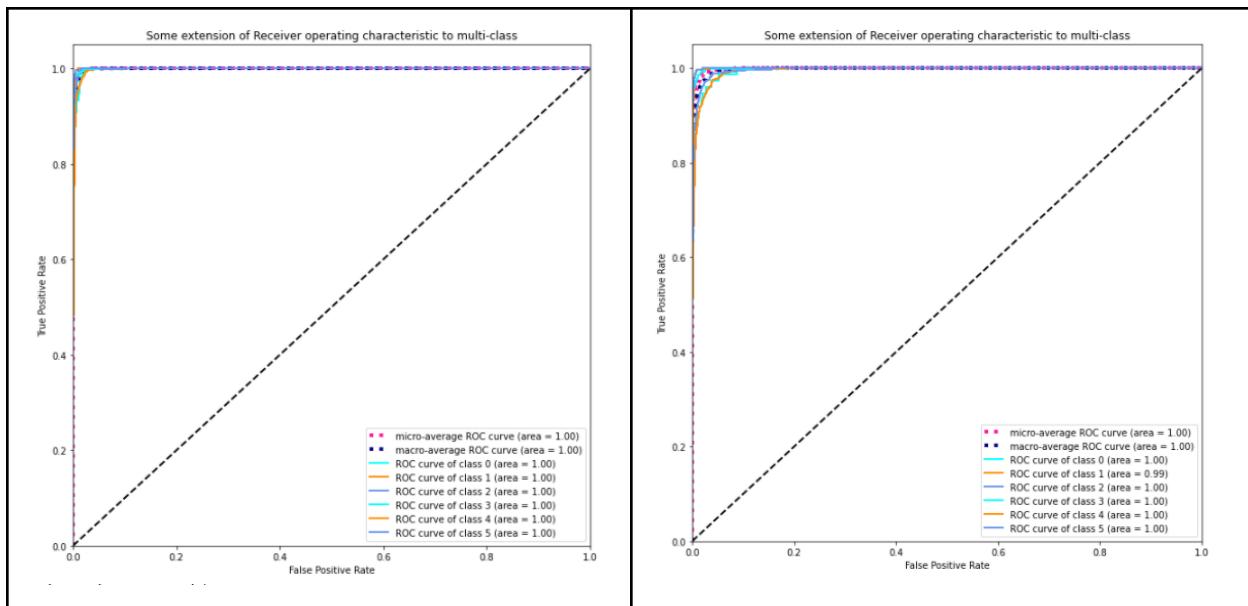
### 1.b) Confusion matrix - Test



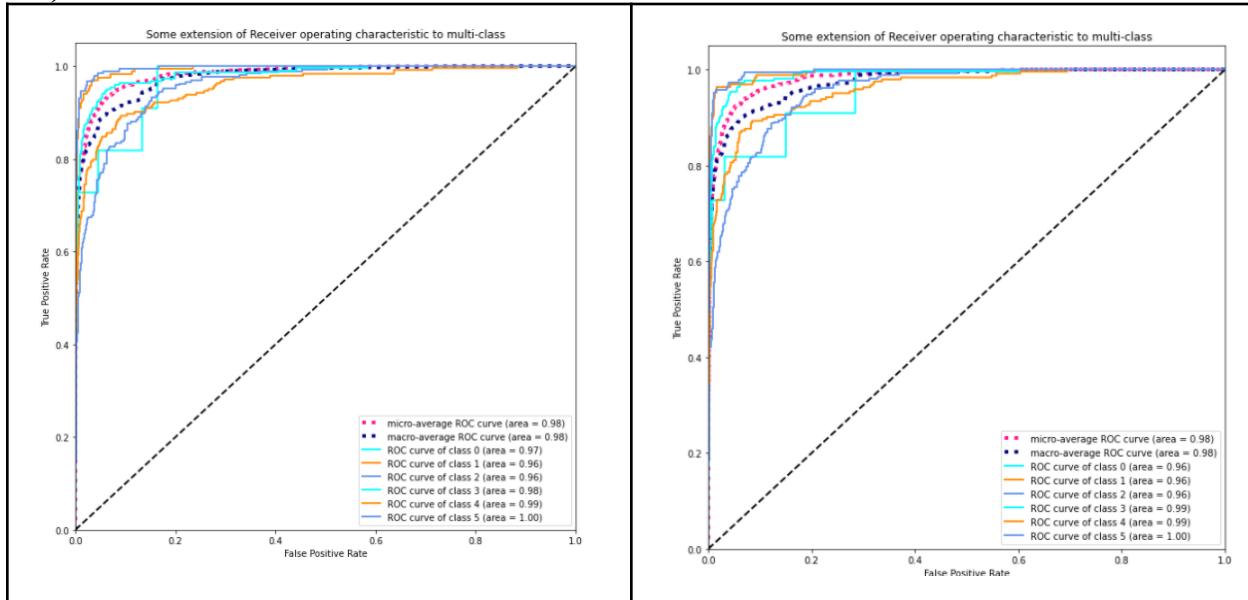


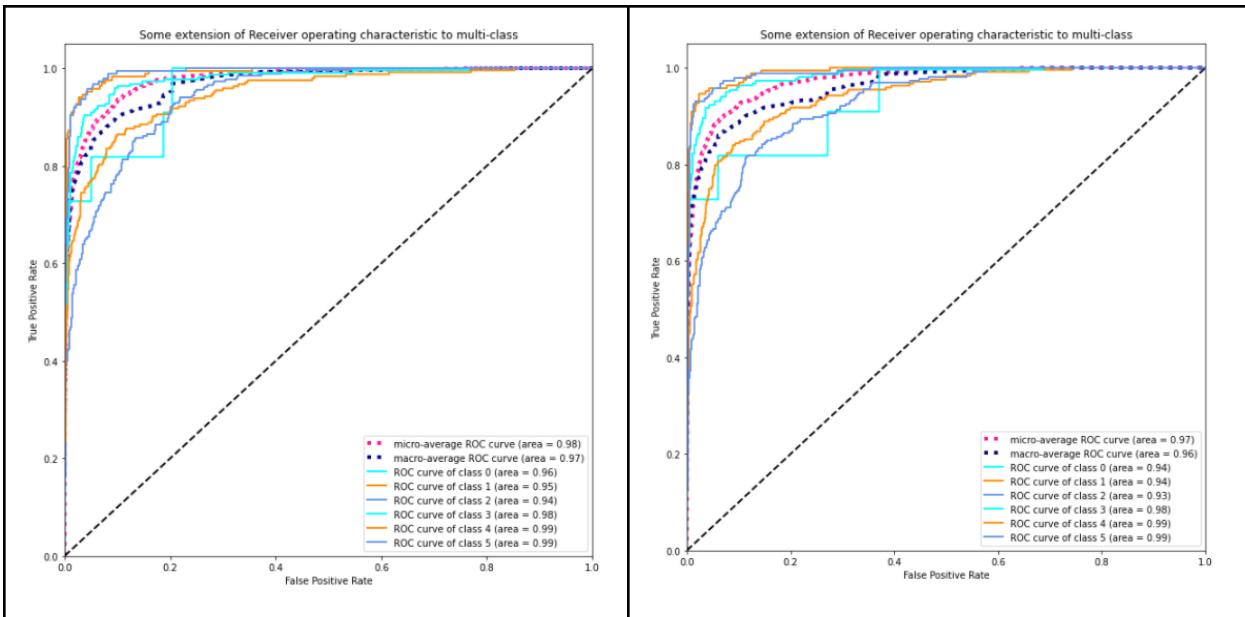
## 2.a) ROC - Train





## 2.b) ROC - Test



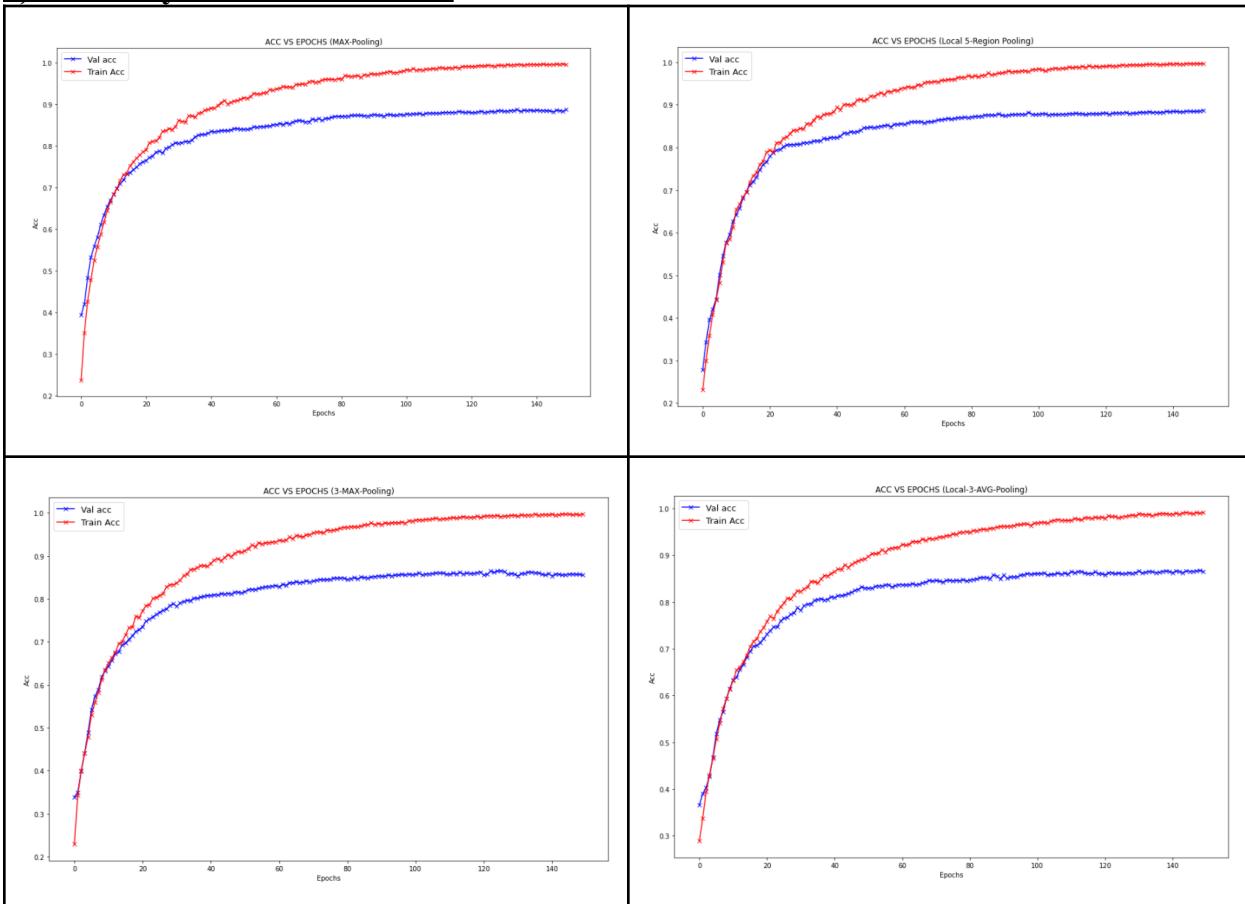


### 3) F1 Measure - train vs test

	precision	recall	f1-score	support		precision	recall	f1-score	support
0	0.857	0.545	0.667	11	0	0.857	0.545	0.667	11
1	0.841	0.852	0.847	243	1	0.838	0.852	0.845	243
2	0.823	0.838	0.831	266	2	0.838	0.816	0.827	266
3	0.909	0.913	0.911	219	3	0.894	0.927	0.910	219
4	0.934	0.940	0.937	166	4	0.946	0.952	0.949	166
5	0.972	0.941	0.956	186	5	0.962	0.952	0.957	186
accuracy			0.886	1091	accuracy			0.887	1091
macro avg	0.889	0.838	0.858	1091	macro avg	0.889	0.841	0.859	1091
weighted avg	0.887	0.886	0.886	1091	weighted avg	0.887	0.887	0.887	1091
	precision	recall	f1-score	support		precision	recall	f1-score	support
0	0.985	0.880	0.930	75	0	1.000	0.893	0.944	75
1	0.986	0.996	0.991	919	1	0.989	1.000	0.995	919
2	0.998	0.993	0.995	984	2	0.999	0.995	0.997	984
3	0.996	0.999	0.998	1004	3	0.998	1.000	0.999	1004
4	0.997	0.996	0.996	669	4	0.999	0.999	0.999	669
5	0.996	0.999	0.997	710	5	1.000	1.000	1.000	710
accuracy			0.994	4361	accuracy			0.997	4361
macro avg	0.993	0.977	0.984	4361	macro avg	0.997	0.981	0.989	4361
weighted avg	0.994	0.994	0.994	4361	weighted avg	0.997	0.997	0.997	4361

	precision	recall	f1-score	support		precision	recall	f1-score	support	
0	0.857	0.545	0.667	11		0	0.857	0.545	0.667	11
1	0.819	0.840	0.829	243		1	0.820	0.827	0.824	243
2	0.820	0.786	0.802	266		2	0.789	0.789	0.789	266
3	0.855	0.918	0.885	219		3	0.849	0.900	0.874	219
4	0.923	0.934	0.928	166		4	0.937	0.898	0.917	166
5	0.955	0.909	0.931	186		5	0.934	0.914	0.924	186
accuracy			0.865	1091	accuracy			0.855	1091	
macro avg	0.871	0.822	0.840	1091	macro avg	0.865	0.812	0.832	1091	
weighted avg	0.866	0.865	0.865	1091	weighted avg	0.856	0.855	0.855	1091	
	precision	recall	f1-score	support		precision	recall	f1-score	support	
0	0.983	0.773	0.866	75		0	1.000	0.880	0.936	75
1	0.974	0.995	0.984	919		1	0.991	0.999	0.995	919
2	0.995	0.988	0.991	984		2	0.999	0.999	0.999	984
3	0.997	0.997	0.997	1004		3	0.999	0.999	0.999	1004
4	0.994	0.999	0.996	669		4	0.997	1.000	0.999	669
5	0.999	1.000	0.999	710		5	1.000	1.000	1.000	710
accuracy			0.991	4361	accuracy			0.997	4361	
macro avg	0.990	0.959	0.972	4361	macro avg	0.998	0.979	0.988	4361	
weighted avg	0.991	0.991	0.991	4361	weighted avg	0.997	0.997	0.997	4361	

#### 4) Accuracy curve train vs test



### **5) Accuracy curve train vs test (Epochs)**

1-max		5-local	
Epoch: 10   Epoch Time: 0m 0s	Train Loss: 1.154   Train Acc: 66.46%	Epoch: 10   Epoch Time: 0m 0s	Train Loss: 1.263   Train Acc: 61.31%
Val. Loss: 1.147   Test. Acc: 66.91%		Val. Loss: 1.251   Test. Acc: 62.69%	
Epoch: 20   Epoch Time: 0m 0s	Train Loss: 0.748   Train Acc: 78.48%	Epoch: 20   Epoch Time: 0m 0s	Train Loss: 0.785   Train Acc: 78.89%
Val. Loss: 0.781   Test. Acc: 76.17%		Val. Loss: 0.819   Test. Acc: 76.63%	
Epoch: 30   Epoch Time: 0m 0s	Train Loss: 0.530   Train Acc: 84.60%	Epoch: 30   Epoch Time: 0m 0s	Train Loss: 0.546   Train Acc: 84.51%
Val. Loss: 0.602   Test. Acc: 80.66%		Val. Loss: 0.610   Test. Acc: 80.75%	
Epoch: 40   Epoch Time: 0m 0s	Train Loss: 0.402   Train Acc: 88.63%	Epoch: 40   Epoch Time: 0m 0s	Train Loss: 0.402   Train Acc: 88.36%
Val. Loss: 0.508   Test. Acc: 82.95%		Val. Loss: 0.507   Test. Acc: 82.31%	
Epoch: 50   Epoch Time: 0m 0s	Train Loss: 0.317   Train Acc: 91.10%	Epoch: 50   Epoch Time: 0m 0s	Train Loss: 0.311   Train Acc: 91.23%
Val. Loss: 0.455   Test. Acc: 83.96%		Val. Loss: 0.449   Test. Acc: 84.60%	
Epoch: 60   Epoch Time: 0m 0s	Train Loss: 0.254   Train Acc: 93.32%	Epoch: 60   Epoch Time: 0m 0s	Train Loss: 0.238   Train Acc: 93.86%
Val. Loss: 0.418   Test. Acc: 84.97%		Val. Loss: 0.413   Test. Acc: 85.52%	
Epoch: 70   Epoch Time: 0m 0s	Train Loss: 0.204   Train Acc: 94.77%	Epoch: 70   Epoch Time: 0m 0s	Train Loss: 0.189   Train Acc: 95.54%
Val. Loss: 0.393   Test. Acc: 85.70%		Val. Loss: 0.387   Test. Acc: 86.16%	
Epoch: 80   Epoch Time: 0m 0s	Train Loss: 0.167   Train Acc: 96.10%	Epoch: 80   Epoch Time: 0m 0s	Train Loss: 0.147   Train Acc: 96.83%
Val. Loss: 0.375   Test. Acc: 86.98%		Val. Loss: 0.370   Test. Acc: 86.98%	
Epoch: 90   Epoch Time: 0m 0s	Train Loss: 0.133   Train Acc: 97.28%	Epoch: 90   Epoch Time: 0m 0s	Train Loss: 0.117   Train Acc: 97.44%
Val. Loss: 0.361   Test. Acc: 87.35%		Val. Loss: 0.358   Test. Acc: 87.53%	
Epoch: 100   Epoch Time: 0m 0s	Train Loss: 0.103   Train Acc: 97.89%	Epoch: 100   Epoch Time: 0m 0s	Train Loss: 0.094   Train Acc: 98.26%
Val. Loss: 0.353   Test. Acc: 87.35%		Val. Loss: 0.349   Test. Acc: 87.72%	
Epoch: 110   Epoch Time: 0m 0s	Train Loss: 0.084   Train Acc: 98.46%	Epoch: 110   Epoch Time: 0m 0s	Train Loss: 0.072   Train Acc: 98.91%
Val. Loss: 0.348   Test. Acc: 87.81%		Val. Loss: 0.344   Test. Acc: 87.90%	
Epoch: 120   Epoch Time: 0m 0s	Train Loss: 0.065   Train Acc: 99.07%	Epoch: 120   Epoch Time: 0m 0s	Train Loss: 0.060   Train Acc: 99.09%
Val. Loss: 0.345   Test. Acc: 88.08%		Val. Loss: 0.342   Test. Acc: 87.99%	
Epoch: 130   Epoch Time: 0m 0s	Train Loss: 0.055   Train Acc: 99.30%	Epoch: 130   Epoch Time: 0m 0s	Train Loss: 0.045   Train Acc: 99.39%
Val. Loss: 0.344   Test. Acc: 88.36%		Val. Loss: 0.341   Test. Acc: 87.99%	
Epoch: 140   Epoch Time: 0m 0s	Train Loss: 0.045   Train Acc: 99.48%	Epoch: 140   Epoch Time: 0m 0s	Train Loss: 0.036   Train Acc: 99.59%
Val. Loss: 0.343   Test. Acc: 88.36%		Val. Loss: 0.342   Test. Acc: 88.36%	
Epoch: 150   Epoch Time: 0m 0s	Train Loss: 0.038   Train Acc: 99.43%	Epoch: 150   Epoch Time: 0m 0s	Train Loss: 0.029   Train Acc: 99.68%
Val. Loss: 0.345   Test. Acc: 88.63%		Val. Loss: 0.344   Test. Acc: 88.73%	
.	.	.	.
3-max		3-local-average	
Epoch: 10   Epoch Time: 0m 0s	Train Loss: 1.196   Train Acc: 63.47%	Epoch: 10   Epoch Time: 0m 0s	Train Loss: 1.220   Train Acc: 61.29%
Val. Loss: 1.201   Test. Acc: 63.24%		Val. Loss: 1.216   Test. Acc: 61.50%	
Epoch: 20   Epoch Time: 0m 0s	Train Loss: 0.778   Train Acc: 75.65%	Epoch: 20   Epoch Time: 0m 0s	Train Loss: 0.814   Train Acc: 74.47%
Val. Loss: 0.832   Test. Acc: 72.87%		Val. Loss: 0.846   Test. Acc: 72.14%	
Epoch: 30   Epoch Time: 0m 0s	Train Loss: 0.552   Train Acc: 83.31%	Epoch: 30   Epoch Time: 0m 0s	Train Loss: 0.573   Train Acc: 82.40%
Val. Loss: 0.645   Test. Acc: 78.83%		Val. Loss: 0.661   Test. Acc: 78.74%	
Epoch: 40   Epoch Time: 0m 0s	Train Loss: 0.420   Train Acc: 87.56%	Epoch: 40   Epoch Time: 0m 0s	Train Loss: 0.449   Train Acc: 86.11%
Val. Loss: 0.550   Test. Acc: 80.75%		Val. Loss: 0.562   Test. Acc: 81.12%	
Epoch: 50   Epoch Time: 0m 0s	Train Loss: 0.319   Train Acc: 90.80%	Epoch: 50   Epoch Time: 0m 0s	Train Loss: 0.351   Train Acc: 89.20%
Val. Loss: 0.494   Test. Acc: 81.39%		Val. Loss: 0.506   Test. Acc: 82.95%	
Epoch: 60   Epoch Time: 0m 0s	Train Loss: 0.251   Train Acc: 93.23%	Epoch: 60   Epoch Time: 0m 0s	Train Loss: 0.284   Train Acc: 91.64%
Val. Loss: 0.458   Test. Acc: 83.13%		Val. Loss: 0.471   Test. Acc: 83.68%	
Epoch: 70   Epoch Time: 0m 0s	Train Loss: 0.198   Train Acc: 94.90%	Epoch: 70   Epoch Time: 0m 0s	Train Loss: 0.232   Train Acc: 93.41%
Val. Loss: 0.434   Test. Acc: 83.87%		Val. Loss: 0.448   Test. Acc: 84.42%	
Epoch: 80   Epoch Time: 0m 0s	Train Loss: 0.159   Train Acc: 96.63%	Epoch: 80   Epoch Time: 0m 0s	Train Loss: 0.188   Train Acc: 95.02%
Val. Loss: 0.418   Test. Acc: 84.78%		Val. Loss: 0.434   Test. Acc: 84.42%	
Epoch: 90   Epoch Time: 0m 0s	Train Loss: 0.122   Train Acc: 97.55%	Epoch: 90   Epoch Time: 0m 0s	Train Loss: 0.154   Train Acc: 96.08%
Val. Loss: 0.409   Test. Acc: 85.33%		Val. Loss: 0.426   Test. Acc: 84.88%	
Epoch: 100   Epoch Time: 0m 0s	Train Loss: 0.096   Train Acc: 98.07%	Epoch: 100   Epoch Time: 0m 0s	Train Loss: 0.129   Train Acc: 96.83%
Val. Loss: 0.402   Test. Acc: 85.61%		Val. Loss: 0.420   Test. Acc: 85.98%	
Epoch: 110   Epoch Time: 0m 0s	Train Loss: 0.074   Train Acc: 98.62%	Epoch: 110   Epoch Time: 0m 0s	Train Loss: 0.103   Train Acc: 97.37%
Val. Loss: 0.400   Test. Acc: 85.79%		Val. Loss: 0.419   Test. Acc: 85.88%	
Epoch: 120   Epoch Time: 0m 0s	Train Loss: 0.062   Train Acc: 98.91%	Epoch: 120   Epoch Time: 0m 0s	Train Loss: 0.087   Train Acc: 98.08%
Val. Loss: 0.403   Test. Acc: 86.25%		Val. Loss: 0.422   Test. Acc: 85.98%	
Epoch: 130   Epoch Time: 0m 0s	Train Loss: 0.046   Train Acc: 99.37%	Epoch: 130   Epoch Time: 0m 0s	Train Loss: 0.069   Train Acc: 98.44%
Val. Loss: 0.406   Test. Acc: 85.79%		Val. Loss: 0.424   Test. Acc: 86.25%	
Epoch: 140   Epoch Time: 0m 0s	Train Loss: 0.039   Train Acc: 99.55%	Epoch: 140   Epoch Time: 0m 0s	Train Loss: 0.059   Train Acc: 98.75%
Val. Loss: 0.412   Test. Acc: 85.88%		Val. Loss: 0.431   Test. Acc: 86.53%	
Epoch: 150   Epoch Time: 0m 0s	Train Loss: 0.029   Train Acc: 99.73%	Epoch: 150   Epoch Time: 0m 0s	Train Loss: 0.049   Train Acc: 99.14%
Val. Loss: 0.419   Test. Acc: 85.52%		Val. Loss: 0.438   Test. Acc: 86.53%	

## **Contribution of Each Member**

### **1. Shivam Sharma**

Implemented part(2) and report

### **2. Akanksha Shrimal**

Implemented part(1) and report

### **3. Vaibhav Goswami**

Implemented part(2) and report